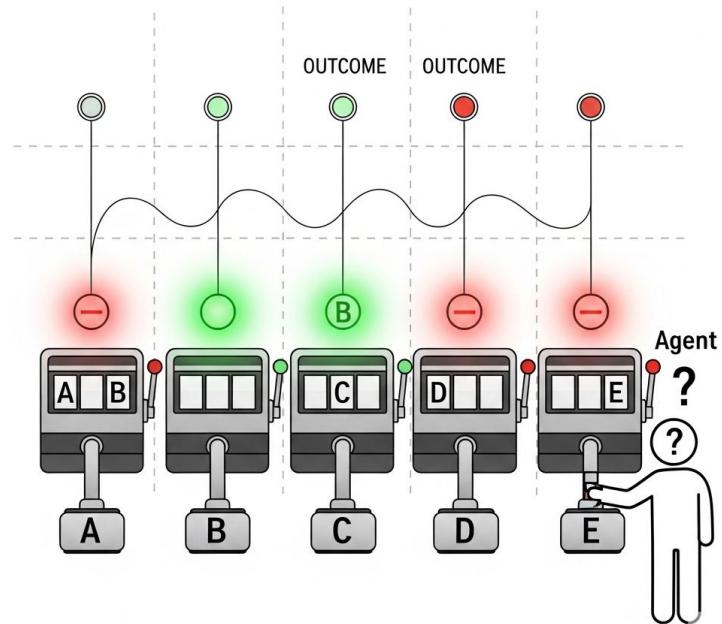


# Multi-Armed bandits



# Agenda

- The Multi-ARmed Bandit (MAB) Problem
- Greedy/Epsilon-Greedy
- Upper Confidence Bound (UCB)
- Thompson Sampling
- Modern Hypothesis Testing

# Multi-Armed Bandit (MAB) Problem

- K Slot Machines  $\{1, 2, \dots, K\}$  (aka "Bandits" with "Arms").



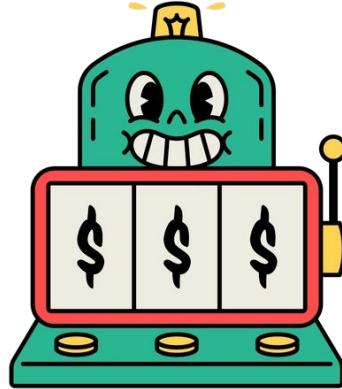
# Multi-Armed Bandit (MAB) Problem

- K Slot Machines  $\{1, 2, \dots, K\}$  (aka "Bandits" with "Arms").
- At each time step  $t=1, 2, \dots, T$ : Pull an arm  $a_t \in \{1, 2, \dots, K\}$  and observe random reward (each arm is independent, and has some reward distribution which doesn't change over time).



# Multi-Armed Bandit (MAB) Problem

- K Slot Machines  $\{1, 2, \dots, K\}$  (aka "Bandits" with "Arms").
- At each time step  $t=1, 2, \dots, T$ : Pull an arm  $a_t \in \{1, 2, \dots, K\}$  and observe random reward (each arm is independent, and has some reward distribution which doesn't change over time).
- Goal: Maximize total (expected) reward after  $T$  time steps.



# Multi-Armed Bandit (MAB) Problem

- K Slot Machines  $\{1, 2, \dots, K\}$  (aka "Bandits" with "Arms").
- At each time step  $t=1, 2, \dots, T$ : Pull an  $a_t \in \{1, 2, \dots, K\}$  and observe random reward (each arm is independent, and has some reward distribution which doesn't change over time).
- Goal: Maximize total (expected) reward after  $T$  time steps.



- Problem: At each time step, decide which arm to pull based on past history of rewards.

# Multi-Armed Bandit (MAB) Problem

Below has the reward distribution of each of the K=3 arms.

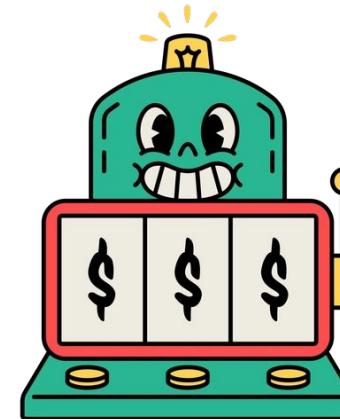
What's your strategy to maximize your total (expected) reward?



$$Poi(\lambda = 1.36)$$

$$Bin(n = 10, p = 0.4)$$

$$\mathcal{N}(\mu = -1, \sigma^2 = 4)$$



# MULTI-ARMED BANDIT (MAB) PROBLEM

Below has the reward distribution of each of the K=3 arms.

What's your strategy to maximize your total (expected) reward?



$Poi(\lambda = 1.36)$



$Bin(n = 10, p = 0.4)$



$\mathcal{N}(\mu = -1, \sigma^2 = 4)$

Pull arm 2 every time since it has the highest expected reward!

# MULTI-ARMED BANDIT (MAB) PROBLEM

Well actually, we don't know the reward distributions :(.



# MULTI-ARMED BANDIT (MAB) PROBLEM

Well actually, we don't know the reward distributions :(.

Have to **estimate** all K expectations



# MULTI-ARMED BANDIT (MAB) PROBLEM

Well actually, we don't know the reward distributions :(.

Have to **estimate** all K expectations, WHILE simultaneously maximizing reward!



# MULTI-ARMED BANDIT (MAB) PROBLEM

Well actually, we don't know the reward distributions :(.

Have to **estimate** all K expectations, WHILE simultaneously maximizing reward!



This is a hard problem - we know nothing about the K reward distributions!

# MULTI-ARMED BANDIT (MAB) PROBLEM

Need to balance the tradeoff between:

**Exploitation:** Pulling arm(s) we know to be "good".

**Exploration:** Pulling other arms in the hopes they are also "good" or even better.



# BERNOULLI BANDITS

We will handle the case of Bernoulli-bandits. That is, reward of arm  $a \in \{1, 2, \dots, n\}$  is  $\text{Ber}(p_a)$ .



$\text{Ber}(p_1)$



$\text{Ber}(p_2)$



$\text{Ber}(p_3)$

# BERNOULLI BANDITS

We will handle the case of Bernoulli-bandits. That is, reward of arm  $a \in \{1, 2, \dots, n\}$  is  $\text{Ber}(p_a)$ .



$\text{Ber}(p_1)$



$\text{Ber}(p_2)$



$\text{Ber}(p_3)$

We don't know these!

# BERNOULLI BANDITS

We will handle the case of Bernoulli-bandits. That is, reward of arm  $a \in \{1, 2, \dots, n\}$  is  $\text{Ber}(p_a)$ .

Observe: The expected reward of arm  $a$   $p_a$  is the expectation of Bernoulli.



$Ber(p_1)$



$Ber(p_2)$



$Ber(p_3)$

We don't know these!

# REGRET

Regret is the difference between:

- The best possible expected reward (always pull the best arm)
- The actual reward you got

# REGRET

Regret is the difference between:

- The best possible expected reward (always pull the best arm)
- The actual reward you got

Let  $p^* = \max_{i \in \{1, 2, \dots, K\}} p_i$  denote the highest expected reward from one of the  $K$  arms.

# REGRET

Regret is the difference between:

- The best possible expected reward (always pull the best arm)
- The actual reward you got

Let  $p^* = \max_{i \in \{1, 2, \dots, K\}} p_i$  denote the highest expected reward from one of the  $K$  arms.

$$\text{Regret}(T) = Tp^* - \text{Reward}(T)$$

At some "time"  $T$  (after  $T$  arm pulls)

# REGRET

Regret is the difference between:

- The best possible expected reward (always pull the best arm)
- The actual reward you got

Let  $p^* = \max_{i \in \{1, 2, \dots, K\}} p_i$  denote the highest expected reward from one of the  $K$  arms.

$$\text{Regret}(T) = Tp^* - \text{Reward}(T)$$

At some "time"  $T$  (after  $T$  arm pulls)

# REGRET

Regret is the difference between:

- The best possible expected reward (always pull the best arm)
- The actual reward you got

Let  $p^* = \max_{i \in \{1, 2, \dots, K\}} p_i$  denote the highest expected reward from one of the  $K$  arms.

$$\text{Regret}(T) = Tp^* - \text{Reward}(T)$$

At some "time"  $T$  (after  $T$  arm pulls)

# REGRET

Regret is the difference between:

- The best possible expected reward (always pull the best arm)
- The actual reward you got

Let  $p^* = \max_{i \in \{1, 2, \dots, K\}} p_i$  denote the highest expected reward from one of the  $K$  arms.

$$\text{Regret}(T) = Tp^* - \text{Reward}(T)$$

$$\text{Avg-Regret}(T) = p^* - \frac{\text{Reward}(T)}{T}$$

# REGRET

Regret is the difference between:

- The best possible expected reward (always pull the best arm)
- The actual reward you got

Let  $p^* = \max_{i \in \{1, 2, \dots, K\}} p_i$  denote the highest expected reward from one of the  $K$  arms.

$$\text{Regret}(T) = Tp^* - \text{Reward}(T)$$

$$\text{Avg-Regret}(T) = p^* - \frac{\text{Reward}(T)}{T}$$

Want  $\text{Avg-Regret}(T) \rightarrow 0$  as  $T \rightarrow \infty$ . Minimizing Regret = Maximizing Reward.

# REGRET

Regret is the difference between:

- The best possible expected reward (always pull the best arm)
- The actual reward you got

Let  $p^* = \max_{i \in \{1, 2, \dots, K\}} p_i$  denote the highest expected reward from one of the  $K$  arms.

$$\text{Regret}(T) = Tp^* - \text{Reward}(T)$$

$$\text{Avg-Regret}(T) = p^* - \frac{\text{Reward}(T)}{T}$$

Want  $\text{Avg-Regret}(T) \rightarrow 0$  as  $T \rightarrow \infty$ . Minimizing Regret = Maximizing Reward.

**Random Quote:** "I'd rather regret the things I've done than regret the things I haven't done." - Anonymous

# (BERNOULLI) BANDIT FRAMEWORK



How do we choose an arm at each time step (depending on past history), to maximize our total reward?

---

## Algorithm 1 (Bernoulli) Bandit Framework

---

1: Have  $K$  arms, where pulling arm  $i \in \{1, \dots, K\}$  gives  $Ber(p_i)$  reward       $\triangleright p_i$ 's all unknown.

# (BERNOULLI) BANDIT FRAMEWORK



How do we choose an arm at each time step (depending on past history), to maximize our total reward?

---

**Algorithm 1** (Bernoulli) Bandit Framework

---

# (BERNOULLI) BANDIT FRAMEWORK



How do we choose an arm at each time step (depending on past history), to maximize our total reward?

---

**Algorithm 1** (Bernoulli) Bandit Framework

# (BERNOULLI) BANDIT FRAMEWORK



How do we choose an arm at each time step (depending on past history), to maximize our total reward?

---

**Algorithm 1** (Bernoulli) Bandit Framework



This is the focus of the rest of this lecture!

# MOTIVATION: CLINICAL TRIALS

K = 4 Arms (Treatments)



# MOTIVATION: CLINICAL TRIALS

$K = 4$  Arms (Treatments)



For patient  $t$ , prescribe treatment  $a_t \in \{1,2,3,4\}$ .

# MOTIVATION: CLINICAL TRIALS

$K = 4$  Arms (Treatments)



For patient  $t$ , prescribe treatment  $a_t$  in  $\{1,2,3,4\}$ .

Observe reward  $r_t$  in  $\{0, 1\}$ . (1 if healed, 0 if not)

# MOTIVATION: CLINICAL TRIALS

$K = 4$  Arms (Treatments)



For patient  $t$ , prescribe treatment  $a_t$  in  $\{1,2,3,4\}$ .

Observe reward  $r_t$  in  $\{0, 1\}$ . (1 if healed, 0 if not)

Maximize: Total number of patients healed.

# MOTIVATION: RECOMMENDING MOVIES



K Movies

For visitor  $t$ , recommend movie  $a_t$  in  $\{1, 2, \dots, K\}$ .

# MOTIVATION: RECOMMENDING MOVIES



K Movies

For visitor  $t$ , recommend movie  $a_t$  in  $\{1, 2, \dots, K\}$ .

Observe reward  $r_t$  in  $\{1, 2, 3, 4, 5\}$ . (rating)

Maximize: Total/average rating of recommendations.

# MOTIVATION: REAL LIFE?? (FOOD)



K Cuisines/Dishes (a ton)

For meal  $t$ , eat dish  $a_t$  in  $\{1, 2, \dots, K\}$ .



# MOTIVATION: REAL LIFE?? (FOOD)



K Cuisines/Dishes (a ton)

For meal  $t$ , eat dish  $a_t$  in  $\{1, 2, \dots, K\}$ .

Observe reward  $r_t$  in  $\{1, 2, 3, 4, 5\}$ . (happiness rating)

Maximize: Total/average happiness :)



# MOTIVATION: REAL LIFE?? (FOOD)



K Cuisines/Dishes (a ton)

For meal  $t$ , eat dish  $a_t$  in  $\{1, 2, \dots, K\}$ .

Observe reward  $r_t$  in  $\{1, 2, 3, 4, 5\}$ . (happiness rating)

Maximize: Total/average happiness :)

The Question of the Day: Explore or Exploit????



# MOTIVATION: REAL LIFE?? (ACTIVITIES)



K Activities

On day  $t$ , do activity  $a_t \in \{1, 2, \dots, K\}$ .



Observe reward  $r_t$  in  $\{1, 2, 3, 4, 5\}$ . (happiness rating)

Maximize: Total/average happiness :)

The Question of the Day: Explore or Exploit????



ANY IDEAS ON WHAT STRATEGY WE CAN USE???





# GREEDY (NAIVE) ALGORITHM

---

**Algorithm 2** Greedy (Naive) Strategy for Bernoulli Bandits

---

- 1: Choose a number of times  $M$  to pull each arm initially, with  $KM \leq T$ .



# GREEDY (NAIVE) ALGORITHM

---

## Algorithm 2 Greedy (Naive) Strategy for Bernoulli Bandits

---

- 1: Choose a number of times  $M$  to pull each arm initially, with  $KM \leq T$ .
- 2: **for**  $i = 1, 2, \dots, K$  **do**
- 3:     Pull arm  $i$   $M$  times, observing iid rewards  $r_{i1}, \dots, r_{iM} \sim Ber(p_i)$ .
- 4:     Estimate  $\hat{p}_i = \frac{\sum_{j=1}^M r_{ij}}{M}$ .



# GREEDY (NAIVE) ALGORITHM

## Algorithm 2 Greedy (Naive) Strategy for Bernoulli Bandits

- 1: Choose a number of times  $M$  to pull each arm initially, with  $KM \leq T$ .
- 2: **for**  $i = 1, 2, \dots, K$  **do**
- 3:     Pull arm  $i$   $M$  times, observing iid rewards  $r_{i1}, \dots, r_{iM} \sim Ber(p_i)$ .
- 4:     Estimate  $\hat{p}_i = \frac{\sum_{j=1}^M r_{ij}}{M}$ .
- 5: Determine best (empirical) arm  $a^* = \arg \max_{i \in \{1, 2, \dots, K\}} \hat{p}_i$ . ► We could be wrong...

# GREEDY (NAIVE) ALGORITHM



**Algorithm 2** Greedy (Naive) Strategy for Bernoulli Bandits

- 1: Choose a number of times  $M$  to pull each arm initially, with  $KM \leq T$ .
  - 2: **for**  $i = 1, 2, \dots, K$  **do**
  - 3:     Pull arm  $i$   $M$  times, observing iid rewards  $r_{i1}, \dots, r_{iM} \sim Ber(p_i)$ .
  - 4:     Estimate  $\hat{p}_i = \frac{\sum_{j=1}^M r_{ij}}{M}$ .
  - 5: Determine best (empirical) arm  $a^* = \arg \max_{i \in \{1, 2, \dots, K\}} \hat{p}_i$ . ▷ We could be wrong...
  - 6: **for**  $t = KM + 1, KM + 2, \dots, T$  **do:**
  - 7:     Pull arm  $a_t = a^*$ . ▷ Pull the same arm for the rest of time.
  - 8:     Receive reward  $r_t \sim Ber(p_{a_t})$ .

# GREEDY (NAIVE) ALGORITHM



**Algorithm 2** Greedy (Naive) Strategy for Bernoulli Bandits



If we make a mistake, we will regret our decision for the rest of time....

Can we not do all of our exploration at the beginning?

$\epsilon$



# EPSILON-GREEDY ALGORITHM

Explore with probability epsilon!

---

### **Algorithm 3** $\epsilon$ -Greedy Strategy for Bernoulli Bandits

---

1: Choose a number of times  $M$  to pull each arm initially, with  $KM \leq T$ .

2: **for**  $i = 1, 2, \dots, K$  **do**

3:     Pull arm  $i$   $M$  times, observing iid rewards  $r_{i1}, \dots, r_{iM} \sim Ber(p_i)$ .

4:     Estimate  $\hat{p}_i = \frac{\sum_{j=1}^M r_{ij}}{M}$ .



$\epsilon$

# EPSILON-GREEDY ALGORITHM

Explore with probability epsilon!

---

### **Algorithm 3** $\epsilon$ -Greedy Strategy for Bernoulli Bandits

---

1: Choose a number of times  $M$  to pull each arm initially, with  $KM \leq T$ .

2: **for**  $i = 1, 2, \dots, K$  **do**

3:     Pull arm  $i$   $M$  times, observing iid rewards  $r_{i1}, \dots, r_{iM} \sim Ber(p_i)$ .

4:     Estimate  $\hat{p}_i = \frac{\sum_{j=1}^M r_{ij}}{M}$ .

5: **for**  $t = KM + 1, KM + 2, \dots, T$  **do:**

6:     **if**  $Ber(\epsilon) == 1$ : **then**

            ► With probability  $\epsilon$ , explore.

7:         Pull arm  $a_t = Unif(1, K)$  (discrete).

            ► Choose a uniformly random arm.

# EPSILON-GREEDY ALGORITHM

ε



# Explore with probability epsilon!

**Algorithm 3**  $\varepsilon$ -Greedy Strategy for Bernoulli Bandits

- ```

1: Choose a number of times  $M$  to pull each arm initially, with  $KM \leq T$ .
2: for  $i = 1, 2, \dots, K$  do
3:   Pull arm  $i$   $M$  times, observing iid rewards  $r_{i1}, \dots, r_{iM} \sim Ber(p_i)$ .
4:   Estimate  $\hat{p}_i = \frac{\sum_{j=1}^M r_{ij}}{M}$ .
5: for  $t = KM + 1, KM + 2, \dots, T$  do:
6:   if  $Ber(\varepsilon) == 1$ : then                                ▷ With probability  $\varepsilon$ , explore.
7:     Pull arm  $a_t = Unif(1, K)$  (discrete).          ▷ Choose a uniformly random arm.
8:   else   ▷ With probability  $1 - \varepsilon$ , exploit.
9:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \hat{p}_i$ . ▷ Choose arm with highest estimated reward.

```



$\epsilon$

# EPSILON-GREEDY ALGORITHM

Explore with probability epsilon!

---

### Algorithm 3 $\epsilon$ -Greedy Strategy for Bernoulli Bandits

---

- 1: Choose a number of times  $M$  to pull each arm initially, with  $KM \leq T$ .
  - 2: **for**  $i = 1, 2, \dots, K$  **do**
  - 3:     Pull arm  $i$   $M$  times, observing iid rewards  $r_{i1}, \dots, r_{iM} \sim Ber(p_i)$ .
  - 4:     Estimate  $\hat{p}_i = \frac{\sum_{j=1}^M r_{ij}}{M}$ .
  - 5: **for**  $t = KM + 1, KM + 2, \dots, T$  **do:**
  - 6:     **if**  $Ber(\epsilon) == 1$ : **then** ► With probability  $\epsilon$ , explore.
  - 7:         Pull arm  $a_t = Unif(1, K)$  (discrete). ► Choose a uniformly random arm.
  - 8:     **else** ► With probability  $1 - \epsilon$ , exploit.
  - 9:         Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \hat{p}_i$ . ► Choose arm with highest estimated reward.
  - 10:         Receive reward  $r_t \sim Ber(p_{a_t})$ .
  - 11:         Update  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).
-



$\epsilon$

# EPSILON-GREEDY ALGORITHM

Explore with probability epsilon!

---

### Algorithm 3 $\epsilon$ -Greedy Strategy for Bernoulli Bandits

---

- 1: Choose a number of times  $M$  to pull each arm initially, with  $KM \leq T$ .
- 2: **for**  $i = 1, 2, \dots, K$  **do**
- 3:     Pull arm  $i$   $M$  times, observing iid rewards  $r_{i1}, \dots, r_{iM} \sim Ber(p_i)$ .
- 4:     Estimate  $\hat{p}_i = \frac{\sum_{j=1}^M r_{ij}}{M}$ .
- 5: **for**  $t = KM + 1, KM + 2, \dots, T$  **do:**
- 6:     **if**  $Ber(\epsilon) == 1$ : **then** ► With probability  $\epsilon$ , explore.
- 7:         Pull arm  $a_t = Unif(1, K)$  (discrete). ► Choose a uniformly random arm.
- 8:     **else** ► With probability  $1 - \epsilon$ , exploit.
- 9:         Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \hat{p}_i$ . ► Choose arm with highest estimated reward.
- 10:         Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 11:         Update  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).

---

Can we explore more "naturally"?

# RANDOM PICTURE

Explore like a  
new-born kitten!





# UPPER CONFIDENCE BOUND (UCB) ALGORITHM

This algorithm constructs confidence intervals for the estimates of each arm, and chooses the arm with the highest **upper** confidence bound (if the confidence interval is  $[a,b]$ , we compare only the value of  $b$ )



# UPPER CONFIDENCE BOUND (UCB) ALGORITHM

This algorithm constructs confidence intervals for the estimates of each arm, and chooses the arm with the highest **upper** confidence bound (if the confidence interval is  $[a,b]$ , we compare only the value of  $b$ )

---

## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

---

- 1: **for**  $i = 1, 2, \dots, K$  **do**
- 2:     Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
- 3:     Estimate  $\hat{p}_i = r_i/1$ .                                              ► Each estimate  $\hat{p}_i$  will initially either be 1 or 0.



# UPPER CONFIDENCE BOUND (UCB) ALGORITHM

This algorithm constructs confidence intervals for the estimates of each arm, and chooses the arm with the highest **upper** confidence bound (if the confidence interval is [a,b], we compare only the value of b)

---

## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

---

- 1: **for**  $i = 1, 2, \dots, K$  **do**
- 2:     Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
- 3:     Estimate  $\hat{p}_i = r_i/1$ .                                              ► Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
- 4: **for**  $t = K + 1, K + 2, \dots, T$  **do**:
- 5:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm  $i$  was pulled before time  $t$ .



# UPPER CONFIDENCE BOUND (UCB) ALGORITHM

This algorithm constructs confidence intervals for the estimates of each arm, and chooses the arm with the highest upper confidence bound (if the confidence interval is  $[a,b]$ , we compare only the value of  $b$ )

**Algorithm 4** UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

- ```

1: for  $i = 1, 2, \dots, K$  do
2:   Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
3:   Estimate  $\hat{p}_i = r_i/1$ .                                ▷ Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
4: for  $t = K + 1, K + 2, \dots, T$  do:
5:   Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm
       $i$  was pulled before time  $t$ .

```

## Point estimate/ Max-likelihood estimate



# UPPER CONFIDENCE BOUND (UCB) ALGORITHM

This algorithm constructs confidence intervals for the estimates of each arm, and chooses the arm with the highest **upper** confidence bound (if the confidence interval is  $[a,b]$ , we compare only the value of  $b$ )

**Algorithm 4** UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

- ```

1: for  $i = 1, 2, \dots, K$  do
2:   Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
3:   Estimate  $\hat{p}_i = r_i/1$ .                                ▷ Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
4: for  $t = K + 1, K + 2, \dots, T$  do:
5:   Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm
       $i$  was pulled before time  $t$ .

```

Point estimate/  
Max-likelihood estimate

Takes the upper part of of the confidence interval.

# UPPER CONFIDENCE BOUND (UCB) ALGORITHM



This algorithm constructs confidence intervals for the estimates of each arm, and chooses the arm with the highest **upper** confidence bound (if the confidence interval is  $[a,b]$ , we compare only the value of  $b$ )

**Algorithm 4** UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

- ```

1: for  $i = 1, 2, \dots, K$  do
2:   Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
3:   Estimate  $\hat{p}_i = r_i/1$ .                                ▷ Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
4: for  $t = K + 1, K + 2, \dots, T$  do:
5:   Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm
       $i$  was pulled before time  $t$ .
6:   Receive reward  $r_t \sim Ber(p_{a_t})$ .

```



# UPPER CONFIDENCE BOUND (UCB) ALGORITHM

This algorithm constructs confidence intervals for the estimates of each arm, and chooses the arm with the highest **upper** confidence bound (if the confidence interval is  $[a,b]$ , we compare only the value of  $b$ )

**Algorithm 4** UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

- ```

1: for  $i = 1, 2, \dots, K$  do
2:   Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
3:   Estimate  $\hat{p}_i = r_i/1$ .                                ▷ Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
4: for  $t = K + 1, K + 2, \dots, T$  do:
5:   Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm
    $i$  was pulled before time  $t$ .
6:   Receive reward  $r_t \sim Ber(p_{a_t})$ .
7:   Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).

```



# UPPER CONFIDENCE BOUND (UCB) ALGORITHM

This algorithm constructs confidence intervals for the estimates of each arm, and chooses the arm with the highest **upper** confidence bound (if the confidence interval is  $[a,b]$ , we compare only the values of  $b$ ).

---

## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

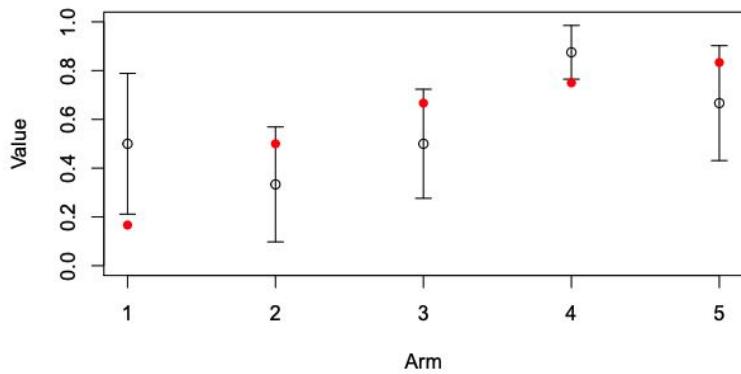
---

- 1: **for**  $i = 1, 2, \dots, K$  **do**
  - 2:     Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
  - 3:     Estimate  $\hat{p}_i = r_i/1$ .                                              ► Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
  - 4: **for**  $t = K + 1, K + 2, \dots, T$  **do**:
  - 5:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm  $i$  was pulled before time  $t$ .
  - 6:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
  - 7:     Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).
- 

Exploration is "baked in": the frequently pulled arms will have narrow confidence intervals (and hence a lower **upper** bound), and the less-frequently pulled arms will have wide intervals (and hence a **higher** upper bound).

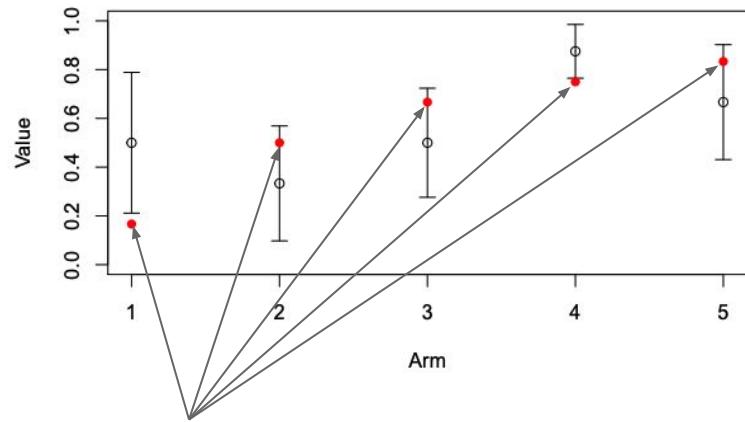
# UCB: CONFIDENCE INTERVALS OVER TIME

Confidence Intervals for Mean of Each Arm: t=10



# UCB: CONFIDENCE INTERVALS OVER TIME

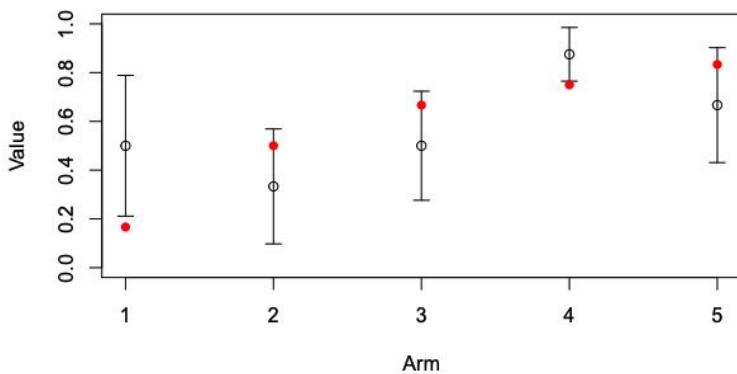
Confidence Intervals for Mean of Each Arm: t=10



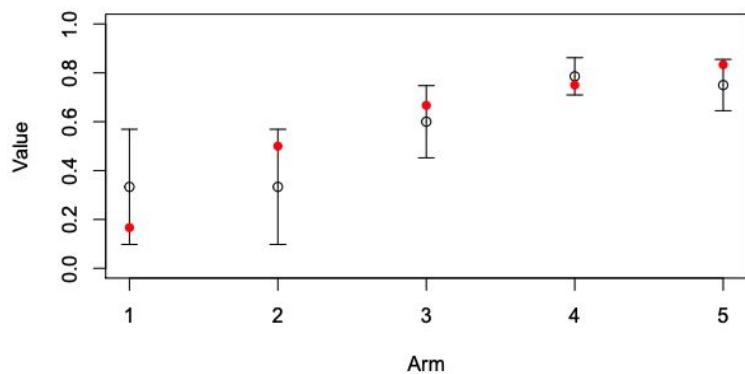
Red Dots: True Means

# UCB: CONFIDENCE INTERVALS OVER TIME

Confidence Intervals for Mean of Each Arm: t=10

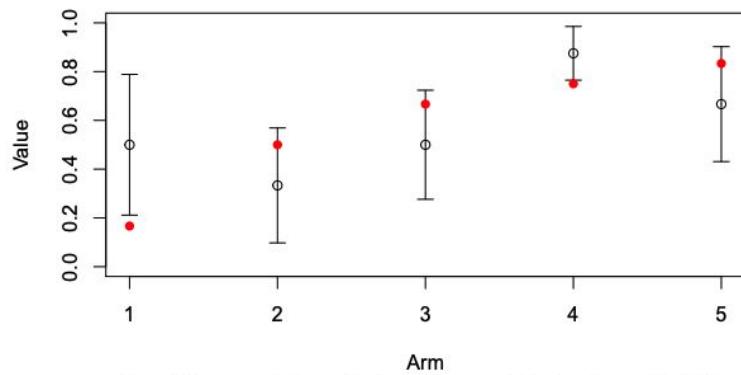


Confidence Intervals for Mean of Each Arm: t=50

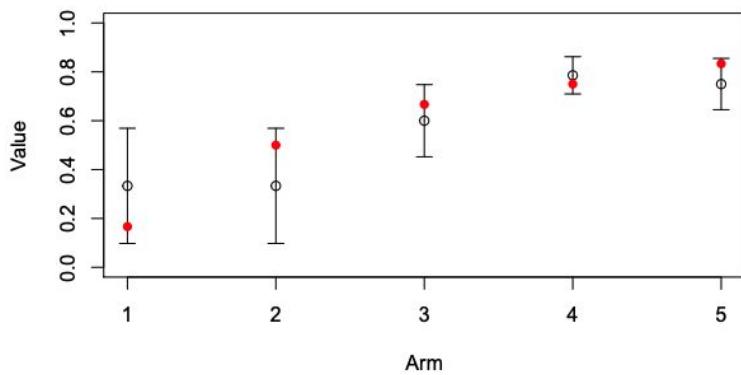


# UCB: CONFIDENCE INTERVALS OVER TIME

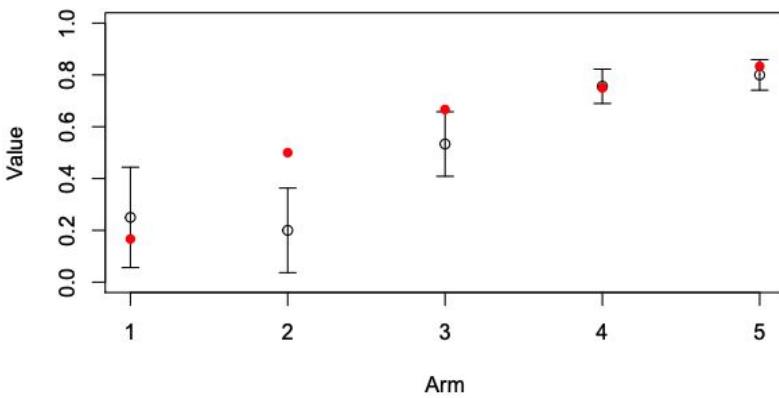
Confidence Intervals for Mean of Each Arm: t=10



Confidence Intervals for Mean of Each Arm: t=50

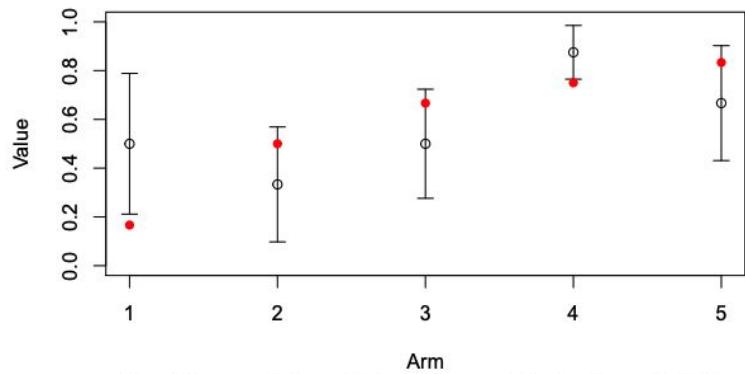


Confidence Intervals for Mean of Each Arm: t=100

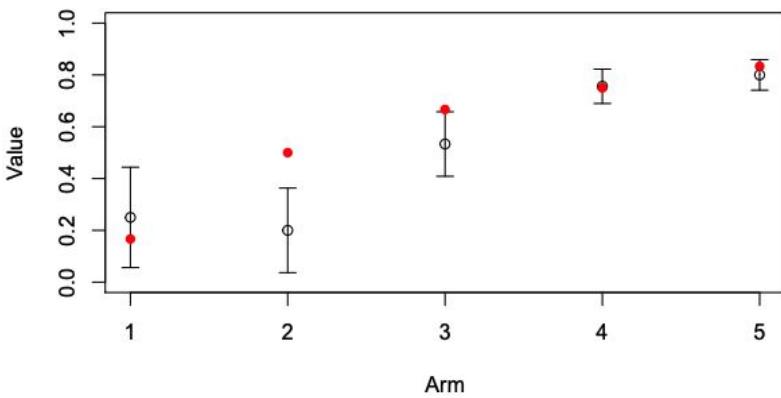


# UCB: CONFIDENCE INTERVALS OVER TIME

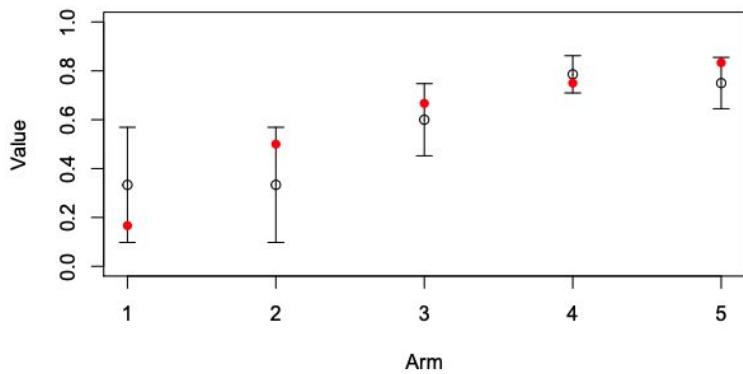
Confidence Intervals for Mean of Each Arm: t=10



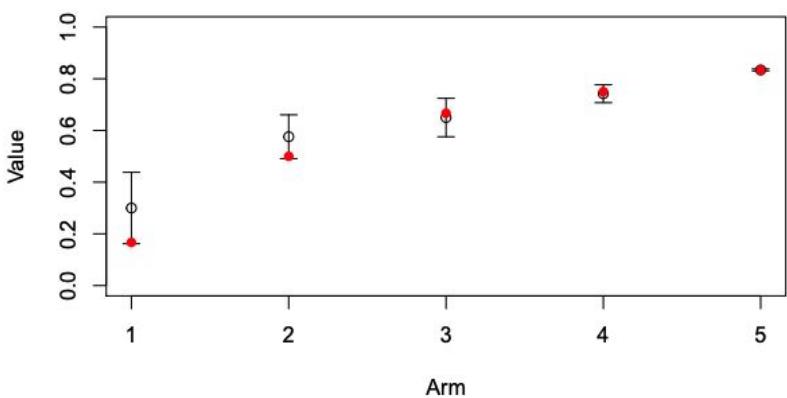
Confidence Intervals for Mean of Each Arm: t=100



Confidence Intervals for Mean of Each Arm: t=50



Confidence Intervals for Mean of Each Arm: t=10000



# UCB EXAMPLE



## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

- 1: **for**  $i = 1, 2, \dots, K$  **do**
- 2:     Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
- 3:     Estimate  $\hat{p}_i = r_i$ . ► Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
- 4: **for**  $t = K + 1, K + 2, \dots, T$  **do**:
- 5:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm  $i$  was pulled before time  $t$ .
- 6:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 7:     Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).

| Arm (i) | True $p_i$ | # Times Pulled | Total Reward | $\hat{p}_i$ | UCB ( $\hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}}$ ) |
|---------|------------|----------------|--------------|-------------|------------------------------------------------------|
| 1       | 0.5        |                |              |             |                                                      |
| 2       | 0.2        |                |              |             |                                                      |
| 3       | 0.9        |                |              |             |                                                      |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
|          |                              |                          |

# UCB EXAMPLE



## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

- 1: **for**  $i = 1, 2, \dots, K$  **do**
- 2:     Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
- 3:     Estimate  $\hat{p}_i = r_i$ . ▷ Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
- 4: **for**  $t = K + 1, K + 2, \dots, T$  **do**:
- 5:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm  $i$  was pulled before time  $t$ .
- 6:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 7:     Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).

| Arm (i) | True $p_i$ | # Times Pulled | Total Reward | $\hat{p}_i$ | UCB ( $\hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}}$ ) |
|---------|------------|----------------|--------------|-------------|------------------------------------------------------|
| 1       | 0.5        |                |              |             |                                                      |
| 2       | 0.2        |                |              |             |                                                      |
| 3       | 0.9        |                |              |             |                                                      |

We don't actually know these...

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
|          |                              |                          |

# UCB EXAMPLE



## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

```

1: for  $i = 1, 2, \dots, K$  do
2:   Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .  
3:   Estimate  $\hat{p}_i = r_i$ . ▷ Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
4: for  $t = K + 1, K + 2, \dots, T$  do:
5:   Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm  $i$  was pulled before time  $t$ .
6:   Receive reward  $r_t \sim Ber(p_{a_t})$ .
7:   Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).
  
```

| Arm (i) | True $p_i$ | # Times Pulled | Total Reward | $\hat{p}_i$ | UCB ( $\hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}}$ ) |
|---------|------------|----------------|--------------|-------------|------------------------------------------------------|
| 1       | 0.5        |                |              |             |                                                      |
| 2       | 0.2        |                |              |             |                                                      |
| 3       | 0.9        |                |              |             |                                                      |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 1        | 1                            | 0                        |

At time 1, we pull arm 1, and observe either a 1 (with probability 0.5) or a 0 (with probability 1-0.5). We happen to observe a 0.

# UCB EXAMPLE



## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

- 1: **for**  $i = 1, 2, \dots, K$  **do**
- 2:     Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
- 3:     Estimate  $\hat{p}_i = r_i$ . ► Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
- 4: **for**  $t = K + 1, K + 2, \dots, T$  **do**:
- 5:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm  $i$  was pulled before time  $t$ .
- 6:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 7:     Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).

| Arm (i) | True $p_i$ | # Times Pulled | Total Reward | $\hat{p}_i$ | UCB ( $\hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}}$ ) |
|---------|------------|----------------|--------------|-------------|------------------------------------------------------|
| 1       | 0.5        | 1              | 0            | 0/1         |                                                      |
| 2       | 0.2        |                |              |             |                                                      |
| 3       | 0.9        |                |              |             |                                                      |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 1        | 1                            | 0                        |

# UCB EXAMPLE



## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

```

1: for  $i = 1, 2, \dots, K$  do
2:   Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .  
3:   Estimate  $\hat{p}_i = r_i$ . ► Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
4: for  $t = K + 1, K + 2, \dots, T$  do:
5:   Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm  $i$  was pulled before time  $t$ .
6:   Receive reward  $r_t \sim Ber(p_{a_t})$ .
7:   Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).
  
```

| Arm (i) | True $p_i$ | # Times Pulled | Total Reward | $\hat{p}_i$ | UCB ( $\hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}}$ ) |
|---------|------------|----------------|--------------|-------------|------------------------------------------------------|
| 1       | 0.5        | 1              | 0            | 0/1         |                                                      |
| 2       | 0.2        |                |              |             |                                                      |
| 3       | 0.9        |                |              |             |                                                      |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 2        | 2                            | 0                        |

At time 2, we pull arm 2, and observe either a 1 (with probability 0.2) or a 0 (with probability 1-0.2). We happen to observe a 0.

# UCB EXAMPLE



## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

- 1: **for**  $i = 1, 2, \dots, K$  **do**
- 2:     Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
- 3:     Estimate  $\hat{p}_i = r_i$ .   ► Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
- 4: **for**  $t = K + 1, K + 2, \dots, T$  **do**:
- 5:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm  $i$  was pulled before time  $t$ .
- 6:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 7:     Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).

| Arm (i) | True $p_i$ | # Times Pulled | Total Reward | $\hat{p}_i$ | UCB ( $\hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}}$ ) |
|---------|------------|----------------|--------------|-------------|------------------------------------------------------|
| 1       | 0.5        | 1              | 0            | 0/1         |                                                      |
| 2       | 0.2        | 1              | 0            | 0/1         |                                                      |
| 3       | 0.9        |                |              |             |                                                      |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 2        | 2                            | 0                        |

# UCB EXAMPLE



## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

```

1: for  $i = 1, 2, \dots, K$  do
2:   Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .  
3:   Estimate  $\hat{p}_i = r_i$ . ▷ Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
4: for  $t = K + 1, K + 2, \dots, T$  do:
5:   Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm  $i$  was pulled before time  $t$ .
6:   Receive reward  $r_t \sim Ber(p_{a_t})$ .
7:   Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).
  
```

| Arm (i) | True $p_i$ | # Times Pulled | Total Reward | $\hat{p}_i$ | UCB ( $\hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}}$ ) |
|---------|------------|----------------|--------------|-------------|------------------------------------------------------|
| 1       | 0.5        | 1              | 0            | 0/1         |                                                      |
| 2       | 0.2        | 1              | 0            | 0/1         |                                                      |
| 3       | 0.9        |                |              |             |                                                      |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 3        | 3                            | 1                        |

At time 3, we pull arm 3, and observe either a 1 (with probability 0.9) or a 0 (with probability 1-0.9). We happen to observe a 1.

# UCB EXAMPLE



## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

- 1: **for**  $i = 1, 2, \dots, K$  **do**
- 2:     Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
- 3:     Estimate  $\hat{p}_i = r_i$ .   ► Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
- 4: **for**  $t = K + 1, K + 2, \dots, T$  **do**:
- 5:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm  $i$  was pulled before time  $t$ .
- 6:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 7:     Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).

| Arm (i) | True $p_i$ | # Times Pulled | Total Reward | $\hat{p}_i$ | UCB ( $\hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}}$ ) |
|---------|------------|----------------|--------------|-------------|------------------------------------------------------|
| 1       | 0.5        | 1              | 0            | 0/1         |                                                      |
| 2       | 0.2        | 1              | 0            | 0/1         |                                                      |
| 3       | 0.9        | 1              | 1            | 1/1         |                                                      |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 3        | 3                            | 1                        |

# UCB EXAMPLE



## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

```

1: for  $i = 1, 2, \dots, K$  do
2:   Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
3:   Estimate  $\hat{p}_i = r_i$ .                                ▷ Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
4: for  $t = K + 1, K + 2, \dots, T$  do:
5:   Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm
    $i$  was pulled before time  $t$ .
6:   Receive reward  $r_t \sim Ber(p_{a_t})$ .
7:   Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).

```

| Arm (i) | True $p_i$ | # Times Pulled | Total Reward | $\hat{p}_i$ | UCB ( $\hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}}$ ) |
|---------|------------|----------------|--------------|-------------|------------------------------------------------------|
| 1       | 0.5        | 1              | 0            | 0/1         |                                                      |
| 2       | 0.2        | 1              | 0            | 0/1         |                                                      |
| 3       | 0.9        | 1              | 1            | 1/1         |                                                      |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 4        |                              |                          |

At time 4, we must compute all our upper confidence bounds, and choose the best one.

# UCB EXAMPLE



## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

- 1: **for**  $i = 1, 2, \dots, K$  **do**
- 2:     Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
- 3:     Estimate  $\hat{p}_i = r_i$ . ► Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
- 4: **for**  $t = K + 1, K + 2, \dots, T$  **do:**
- 5:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm  $i$  was pulled before time  $t$ .
- 6:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 7:     Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).

| Arm (i) | True $p_i$ | # Times Pulled | Total Reward | $\hat{p}_i$ | UCB ( $\hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}}$ ) |
|---------|------------|----------------|--------------|-------------|------------------------------------------------------|
| 1       | 0.5        | 1              | 0            | 0/1         | 1.665                                                |
| 2       | 0.2        | 1              | 0            | 0/1         |                                                      |
| 3       | 0.9        | 1              | 1            | 1/1         |                                                      |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 4        |                              |                          |

$$0 + \sqrt{\frac{2 \ln(4)}{1}} \approx 1.665$$

At time 4, we must compute all our upper confidence bounds, and choose the best one.

# UCB EXAMPLE



## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

- 1: **for**  $i = 1, 2, \dots, K$  **do**
- 2:     Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
- 3:     Estimate  $\hat{p}_i = r_i$ . ► Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
- 4: **for**  $t = K + 1, K + 2, \dots, T$  **do:**
- 5:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm  $i$  was pulled before time  $t$ .
- 6:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 7:     Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).

| Arm (i) | True $p_i$ | # Times Pulled | Total Reward | $\hat{p}_i$ | UCB ( $\hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}}$ ) |
|---------|------------|----------------|--------------|-------------|------------------------------------------------------|
| 1       | 0.5        | 1              | 0            | 0/1         | 1.665                                                |
| 2       | 0.2        | 1              | 0            | 0/1         | 1.665                                                |
| 3       | 0.9        | 1              | 1            | 1/1         |                                                      |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 4        |                              |                          |

$$0 + \sqrt{\frac{2 \ln(4)}{1}} \approx 1.665$$

At time 4, we must compute all our upper confidence bounds, and choose the best one.

# UCB EXAMPLE



## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

- 1: **for**  $i = 1, 2, \dots, K$  **do**
- 2:     Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
- 3:     Estimate  $\hat{p}_i = r_i$ . ► Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
- 4: **for**  $t = K + 1, K + 2, \dots, T$  **do:**
- 5:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm  $i$  was pulled before time  $t$ .
- 6:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 7:     Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).

| Arm (i) | True $p_i$ | # Times Pulled | Total Reward | $\hat{p}_i$ | UCB ( $\hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}}$ ) |
|---------|------------|----------------|--------------|-------------|------------------------------------------------------|
| 1       | 0.5        | 1              | 0            | 0/1         | 1.665                                                |
| 2       | 0.2        | 1              | 0            | 0/1         | 1.665                                                |
| 3       | 0.9        | 1              | 1            | 1/1         | 2.665                                                |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 4        |                              |                          |

$$1 + \sqrt{\frac{2 \ln(4)}{1}} \approx 2.665$$

At time 4, we must compute all our upper confidence bounds, and choose the best one.

# UCB EXAMPLE



## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

- 1: **for**  $i = 1, 2, \dots, K$  **do**
- 2:     Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
- 3:     Estimate  $\hat{p}_i = r_i$ . ► Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
- 4: **for**  $t = K + 1, K + 2, \dots, T$  **do:**
- 5:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm  $i$  was pulled before time  $t$ .
- 6:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 7:     Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).

| Arm (i) | True $p_i$ | # Times Pulled | Total Reward | $\hat{p}_i$ | UCB ( $\hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}}$ ) |
|---------|------------|----------------|--------------|-------------|------------------------------------------------------|
| 1       | 0.5        | 1              | 0            | 0/1         | 1.665                                                |
| 2       | 0.2        | 1              | 0            | 0/1         | 1.665                                                |
| 3       | 0.9        | 1              | 1            | 1/1         | 2.665                                                |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 4        | 3                            |                          |

At time 4, arm 3 has the highest UCB so we pull it.

# UCB EXAMPLE



## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

- 1: **for**  $i = 1, 2, \dots, K$  **do**
- 2:     Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
- 3:     Estimate  $\hat{p}_i = r_i$ . ► Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
- 4: **for**  $t = K + 1, K + 2, \dots, T$  **do**:
- 5:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm  $i$  was pulled before time  $t$ .
- 6:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 7:     Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).

| Arm (i) | True $p_i$ | # Times Pulled | Total Reward | $\hat{p}_i$ | UCB ( $\hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}}$ ) |
|---------|------------|----------------|--------------|-------------|------------------------------------------------------|
| 1       | 0.5        | 1              | 0            | 0/1         | 1.665                                                |
| 2       | 0.2        | 1              | 0            | 0/1         | 1.665                                                |
| 3       | 0.9        | 1              | 1            | 1/1         | 2.665                                                |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 4        | 3                            | 0                        |

At time 4, arm 3 has the highest UCB so we pull it. We observe a reward of 0.

# UCB EXAMPLE



## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

- 1: **for**  $i = 1, 2, \dots, K$  **do**
- 2:     Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
- 3:     Estimate  $\hat{p}_i = r_i$ . ► Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
- 4: **for**  $t = K + 1, K + 2, \dots, T$  **do**:
- 5:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm  $i$  was pulled before time  $t$ .
- 6:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 7:     Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).

| Arm (i) | True $p_i$ | # Times Pulled | Total Reward | $\hat{p}_i$ | UCB ( $\hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}}$ ) |
|---------|------------|----------------|--------------|-------------|------------------------------------------------------|
| 1       | 0.5        | 1              | 0            | 0/1         |                                                      |
| 2       | 0.2        | 1              | 0            | 0/1         |                                                      |
| 3       | 0.9        | 2              | 1            | 1/2         |                                                      |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 4        | 3                            | 0                        |

At time 4, arm 3 has the highest UCB so we pull it. We observe a reward of 0. Then we update our estimate for  $p_3$ .

# UCB EXAMPLE



## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

```

1: for  $i = 1, 2, \dots, K$  do
2:   Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
3:   Estimate  $\hat{p}_i = r_i$ .                                 $\triangleright$  Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
4: for  $t = K + 1, K + 2, \dots, T$  do:
5:   Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm
    $i$  was pulled before time  $t$ .
6:   Receive reward  $r_t \sim Ber(p_{a_t})$ .
7:   Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).

```

| Arm (i) | True $p_i$ | # Times Pulled | Total Reward | $\hat{p}_i$ | UCB ( $\hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}}$ ) |
|---------|------------|----------------|--------------|-------------|------------------------------------------------------|
| 1       | 0.5        | 1              | 0            | 0/1         |                                                      |
| 2       | 0.2        | 1              | 0            | 0/1         |                                                      |
| 3       | 0.9        | 2              | 1            | 1/2         |                                                      |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 5        |                              |                          |

At time 5, we must compute all our upper confidence bounds, and choose the best one.

# UCB EXAMPLE



## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

- 1: **for**  $i = 1, 2, \dots, K$  **do**
- 2:     Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
- 3:     Estimate  $\hat{p}_i = r_i$ . ► Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
- 4: **for**  $t = K + 1, K + 2, \dots, T$  **do:**
- 5:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm  $i$  was pulled before time  $t$ .
- 6:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 7:     Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).

| Arm (i) | True $p_i$ | # Times Pulled | Total Reward | $\hat{p}_i$ | UCB ( $\hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}}$ ) |
|---------|------------|----------------|--------------|-------------|------------------------------------------------------|
| 1       | 0.5        | 1              | 0            | 0/1         | 1.794                                                |
| 2       | 0.2        | 1              | 0            | 0/1         |                                                      |
| 3       | 0.9        | 2              | 1            | 1/2         |                                                      |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 5        |                              |                          |

At time 5, we must compute all our upper confidence bounds, and choose the best one.

$$0 + \sqrt{\frac{2 \ln(5)}{1}} \approx 1.794$$

# UCB EXAMPLE



## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

```

1: for  $i = 1, 2, \dots, K$  do
2:   Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
3:   Estimate  $\hat{p}_i = r_i$ .                                ▷ Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
4: for  $t = K + 1, K + 2, \dots, T$  do:
5:   Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm
    $i$  was pulled before time  $t$ .
6:   Receive reward  $r_t \sim Ber(p_{a_t})$ .
7:   Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).

```

| Arm (i) | True $p_i$ | # Times Pulled | Total Reward | $\hat{p}_i$ | UCB ( $\hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}}$ ) |
|---------|------------|----------------|--------------|-------------|------------------------------------------------------|
| 1       | 0.5        | 1              | 0            | 0/1         | 1.794                                                |
| 2       | 0.2        | 1              | 0            | 0/1         | 1.794                                                |
| 3       | 0.9        | 2              | 1            | 1/2         |                                                      |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 5        |                              |                          |

At time 5, we must compute all our upper confidence bounds, and choose the best one.

$$0 + \sqrt{\frac{2 \ln(5)}{1}} \approx 1.794$$

# UCB EXAMPLE



## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

```

1: for  $i = 1, 2, \dots, K$  do
2:   Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
3:   Estimate  $\hat{p}_i = r_i$ .                                ▷ Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
4: for  $t = K + 1, K + 2, \dots, T$  do:
5:   Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm
    $i$  was pulled before time  $t$ .
6:   Receive reward  $r_t \sim Ber(p_{a_t})$ .
7:   Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).

```

| Arm (i) | True $p_i$ | # Times Pulled | Total Reward | $\hat{p}_i$ | UCB ( $\hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}}$ ) |
|---------|------------|----------------|--------------|-------------|------------------------------------------------------|
| 1       | 0.5        | 1              | 0            | 0/1         | 1.794                                                |
| 2       | 0.2        | 1              | 0            | 0/1         | 1.794                                                |
| 3       | 0.9        | 2              | 1            | 1/2         | 1.769                                                |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 5        |                              |                          |

At time 5, we must compute all our upper confidence bounds, and choose the best one.

$$\frac{1}{2} + \sqrt{\frac{2 \ln(5)}{2}} \approx 1.769$$

# UCB EXAMPLE



## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

- 1: **for**  $i = 1, 2, \dots, K$  **do**
- 2:     Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
- 3:     Estimate  $\hat{p}_i = r_i$ . ► Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
- 4: **for**  $t = K + 1, K + 2, \dots, T$  **do:**
- 5:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm  $i$  was pulled before time  $t$ .
- 6:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 7:     Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).

| Arm (i) | True $p_i$ | # Times Pulled | Total Reward | $\hat{p}_i$ | UCB ( $\hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}}$ ) |
|---------|------------|----------------|--------------|-------------|------------------------------------------------------|
| 1       | 0.5        | 1              | 0            | 0/1         | 1.794                                                |
| 2       | 0.2        | 1              | 0            | 0/1         | 1.794                                                |
| 3       | 0.9        | 2              | 1            | 1/2         | 1.769                                                |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 5        | 1                            |                          |

At time 5, arms 1 and 2 have the highest UCB so we pull one of them (let's break ties by choosing the smaller index arm). So we pull arm 1.

# UCB EXAMPLE



## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

- 1: **for**  $i = 1, 2, \dots, K$  **do**
- 2:     Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
- 3:     Estimate  $\hat{p}_i = r_i$ . ▷ Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
- 4: **for**  $t = K + 1, K + 2, \dots, T$  **do**:
- 5:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm  $i$  was pulled before time  $t$ .
- 6:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 7:     Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).

| Arm (i) | True $p_i$ | # Times Pulled | Total Reward | $\hat{p}_i$ | UCB ( $\hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}}$ ) |
|---------|------------|----------------|--------------|-------------|------------------------------------------------------|
| 1       | 0.5        | 1              | 0            | 0/1         | 1.794                                                |
| 2       | 0.2        | 1              | 0            | 0/1         | 1.794                                                |
| 3       | 0.9        | 2              | 1            | 1/2         | 1.769                                                |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 5        | 1                            | 0                        |

At time 5, arms 1 and 2 have the highest UCB so we pull one of them (let's break ties by choosing the smaller index arm). So we pull arm 1. We observe a reward of 0.

# UCB EXAMPLE



## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

- 1: **for**  $i = 1, 2, \dots, K$  **do**
- 2:     Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
- 3:     Estimate  $\hat{p}_i = r_i$ . ► Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
- 4: **for**  $t = K + 1, K + 2, \dots, T$  **do:**
- 5:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm  $i$  was pulled before time  $t$ .
- 6:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 7:     Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).

| Arm (i) | True $p_i$ | # Times Pulled | Total Reward | $\hat{p}_i$ | UCB ( $\hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}}$ ) |
|---------|------------|----------------|--------------|-------------|------------------------------------------------------|
| 1       | 0.5        | 2              | 0            | 0/2         |                                                      |
| 2       | 0.2        | 1              | 0            | 0/1         |                                                      |
| 3       | 0.9        | 2              | 1            | 1/2         |                                                      |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 5        | 1                            | 0                        |

At time 5, arms 1 and 2 have the highest UCB so we pull one of them (let's break ties by choosing the smaller index arm). So we pull arm 1.

We observe a reward of 0. Then we update our estimate for  $p_1$ .

# UCB EXAMPLE



## Algorithm 4 UCB1 Algorithm (Upper Confidence Bound) for Bernoulli Bandits

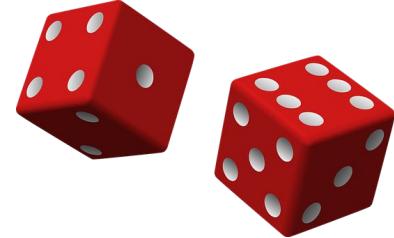
- 1: **for**  $i = 1, 2, \dots, K$  **do**
- 2:     Pull arm  $i$  once, observing  $r_i \sim Ber(p_i)$ .
- 3:     Estimate  $\hat{p}_i = r_i$ . ► Each estimate  $\hat{p}_i$  will initially either be 1 or 0.
- 4: **for**  $t = K + 1, K + 2, \dots, T$  **do:**
- 5:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} \left( \hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}} \right)$ , where  $N_t(i)$  is the number of times arm  $i$  was pulled before time  $t$ .
- 6:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 7:     Update  $N_t(a_t)$  and  $\hat{p}_{a_t}$  (using newly observed reward  $r_t$ ).

| Arm (i) | True $p_i$ | # Times Pulled | Total Reward | $\hat{p}_i$ | UCB ( $\hat{p}_i + \sqrt{\frac{2 \ln(t)}{N_t(i)}}$ ) |
|---------|------------|----------------|--------------|-------------|------------------------------------------------------|
| 1       | 0.5        | 2              | 0            | 0/2         |                                                      |
| 2       | 0.2        | 1              | 0            | 0/1         |                                                      |
| 3       | 0.9        | 2              | 1            | 1/2         |                                                      |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 6        |                              |                          |

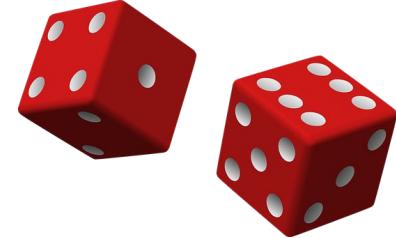
And so on!!! Notice how we started exploring since the confidence bound grows with  $t$  for even the unexplored arms!

# THOMPSON SAMPLING ALGORITHM



Use MAP: Assume a  $\text{Beta}(1,1)$  (Uniform) prior on each unknown probability of reward.

# THOMPSON SAMPLING ALGORITHM



Use MAP: Assume a  $\text{Beta}(1,1)$  (Uniform) prior on each unknown probability of reward.

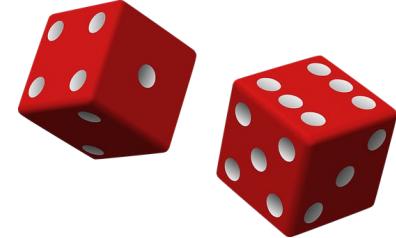
---

## Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

---

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ► Set  $\text{Beta}(\alpha_i, \beta_i)$  prior for each  $p_i$ .

# THOMPSON SAMPLING ALGORITHM



Use MAP: Assume a  $\text{Beta}(1,1)$  (Uniform) prior on each unknown probability of reward.

---

## Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

---

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ▷ Set  $\text{Beta}(\alpha_i, \beta_i)$  prior for each  $p_i$ .
- 2: **for**  $t = 1, 2, \dots, T$  **do**:
- 3:     For each arm  $i$ , get sample  $s_{i,t} \sim \text{Beta}(\alpha_i, \beta_i)$ . ▷ Each is a float in  $[0, 1]$ .

# THOMPSON SAMPLING ALGORITHM



Use MAP: Assume a  $\text{Beta}(1,1)$  (Uniform) prior on each unknown probability of reward.

---

## Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

---

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ▷ Set  $\text{Beta}(\alpha_i, \beta_i)$  prior for each  $p_i$ .
- 2: **for**  $t = 1, 2, \dots, T$  **do:**
- 3:     For each arm  $i$ , get sample  $s_{i,t} \sim \text{Beta}(\alpha_i, \beta_i)$ . ▷ Each is a float in  $[0, 1]$ .
- 4:     Pull arm  $a_t = \arg \max_{i \in \{1,2,\dots,K\}} s_{i,t}$ . ▷ This “bakes in” exploration!

# THOMPSON SAMPLING ALGORITHM



Use MAP: Assume a  $\text{Beta}(1,1)$  (Uniform) prior on each unknown probability of reward.

---

## Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

---

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ▷ Set  $\text{Beta}(\alpha_i, \beta_i)$  prior for each  $p_i$ .
- 2: **for**  $t = 1, 2, \dots, T$  **do:**
- 3:     For each arm  $i$ , get sample  $s_{i,t} \sim \text{Beta}(\alpha_i, \beta_i)$ . ▷ Each is a float in  $[0, 1]$ .
- 4:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} s_{i,t}$ . ▷ This “bakes in” exploration!
- 5:     Receive reward  $r_t \sim \text{Ber}(p_{a_t})$ .

# THOMPSON SAMPLING ALGORITHM



Use MAP: Assume a  $\text{Beta}(1,1)$  (Uniform) prior on each unknown probability of reward.

---

## Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

---

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ▷ Set  $\text{Beta}(\alpha_i, \beta_i)$  prior for each  $p_i$ .
  - 2: **for**  $t = 1, 2, \dots, T$  **do**:
  - 3:   For each arm  $i$ , get sample  $s_{i,t} \sim \text{Beta}(\alpha_i, \beta_i)$ . ▷ Each is a float in  $[0, 1]$ .
  - 4:   Pull arm  $a_t = \arg \max_{i \in \{1,2,\dots,K\}} s_{i,t}$ . ▷ This “bakes in” exploration!
  - 5:   Receive reward  $r_t \sim \text{Ber}(p_{a_t})$ .
  - 6:   **if**  $r_t == 1$  **then**  $\alpha_{a_t} \leftarrow \alpha_{a_t} + 1$ . ▷ Increment number of “successes”.
  - 7:   **else if**  $r_t == 0$  **then**  $\beta_{a_t} \leftarrow \beta_{a_t} + 1$ . ▷ Increment number of “failures”.
-

# THOMPSON SAMPLING ALGORITHM



Use MAP: Assume a  $\text{Beta}(1,1)$  (Uniform) prior on each unknown probability of reward.

---

## Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

---

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ► Set  $\text{Beta}(\alpha_i, \beta_i)$  prior for each  $p_i$ .
  - 2: **for**  $t = 1, 2, \dots, T$  **do:**
  - 3:   For each arm  $i$ , get sample  $s_{i,t} \sim \text{Beta}(\alpha_i, \beta_i)$ . ► Each is a float in  $[0, 1]$ .
  - 4:   Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} s_{i,t}$ . ► This “bakes in” exploration!
  - 5:   Receive reward  $r_t \sim \text{Ber}(p_{a_t})$ .
  - 6:   **if**  $r_t == 1$  **then**  $\alpha_{a_t} \leftarrow \alpha_{a_t} + 1$ . ► Increment number of “successes”.
  - 7:   **else if**  $r_t == 0$  **then**  $\beta_{a_t} \leftarrow \beta_{a_t} + 1$ . ► Increment number of “failures”.
- 

The exploration comes in since we sample from each Beta distribution, rather than just choosing the one with largest expectation or mode (greedy).

# THOMPSON

## EXAMPLE



### Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ▷ Set  $Beta(\alpha_i, \beta_i)$  prior for each  $p_i$ .
- 2: **for**  $t = 1, 2, \dots, T$  **do**:
- 3:     For each arm  $i$ , get sample  $s_{i,t} \sim Beta(\alpha_i, \beta_i)$ . ▷ Each is a float in  $[0, 1]$ .
- 4:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} s_{i,t}$ . ▷ This “bakes in” exploration!
- 5:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 6:     **if**  $r_t == 1$  **then**  $\alpha_{a_t} \leftarrow \alpha_{a_t} + 1$ . ▷ Increment number of “successes”.
- 7:     **else if**  $r_t == 0$  **then**  $\beta_{a_t} \leftarrow \beta_{a_t} + 1$ . ▷ Increment number of “failures”.

| Arm (i) | True $p_i$ | $\alpha_i$ | $\beta_i$ | $s_{i,t}$ |
|---------|------------|------------|-----------|-----------|
| 1       | 0.5        |            |           |           |
| 2       | 0.2        |            |           |           |
| 3       | 0.9        |            |           |           |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
|          |                              |                          |

# THOMPSON

## EXAMPLE



### Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ▷ Set  $Beta(\alpha_i, \beta_i)$  prior for each  $p_i$ .
- 2: **for**  $t = 1, 2, \dots, T$  **do**:
- 3:     For each arm  $i$ , get sample  $s_{i,t} \sim Beta(\alpha_i, \beta_i)$ . ▷ Each is a float in  $[0, 1]$ .
- 4:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} s_{i,t}$ . ▷ This “bakes in” exploration!
- 5:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 6:     **if**  $r_t == 1$  **then**  $\alpha_{a_t} \leftarrow \alpha_{a_t} + 1$ . ▷ Increment number of “successes”.
- 7:     **else if**  $r_t == 0$  **then**  $\beta_{a_t} \leftarrow \beta_{a_t} + 1$ . ▷ Increment number of “failures”.

| Arm (i) | True $p_i$ | $\alpha_i$ | $\beta_i$ | $s_{i,t}$ |
|---------|------------|------------|-----------|-----------|
| 1       | 0.5        | 1          | 1         |           |
| 2       | 0.2        | 1          | 1         |           |
| 3       | 0.9        | 1          | 1         |           |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 1        |                              |                          |

# THOMPSON

## EXAMPLE

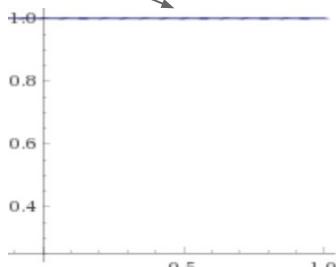


### Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ▷ Set  $Beta(\alpha_i, \beta_i)$  prior for each  $p_i$ .
- 2: **for**  $t = 1, 2, \dots, T$  **do**:
- 3:     For each arm  $i$ , get sample  $s_{i,t} \sim Beta(\alpha_i, \beta_i)$ . ▷ Each is a float in  $[0, 1]$ .
- 4:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} s_{i,t}$ . ▷ This “bakes in” exploration!
- 5:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 6:     **if**  $r_t == 1$  **then**  $\alpha_{a_t} \leftarrow \alpha_{a_t} + 1$ . ▷ Increment number of “successes”.
- 7:     **else if**  $r_t == 0$  **then**  $\beta_{a_t} \leftarrow \beta_{a_t} + 1$ . ▷ Increment number of “failures”.

| Arm (i) | True $p_i$ | $\alpha_i$ | $\beta_i$ | $s_{i,t}$ |
|---------|------------|------------|-----------|-----------|
| 1       | 0.5        | 1          | 1         | 0.43      |
| 2       | 0.2        | 1          | 1         |           |
| 3       | 0.9        | 1          | 1         |           |

Sample from  $Beta(1,1)$  density →



| Time (t) | Arm Pulled ( $a_t$ ) | Reward ( $r_t$ ) |
|----------|----------------------|------------------|
| 1        |                      |                  |

# THOMPSON

## EXAMPLE

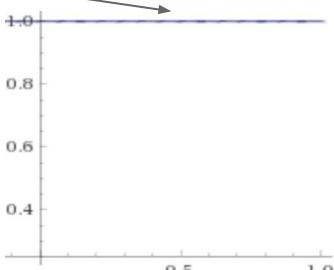


### Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ▷ Set  $Beta(\alpha_i, \beta_i)$  prior for each  $p_i$ .
- 2: **for**  $t = 1, 2, \dots, T$  **do**:
- 3:     For each arm  $i$ , get sample  $s_{i,t} \sim Beta(\alpha_i, \beta_i)$ . ▷ Each is a float in  $[0, 1]$ .
- 4:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} s_{i,t}$ . ▷ This “bakes in” exploration!
- 5:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 6:     **if**  $r_t == 1$  **then**  $\alpha_{a_t} \leftarrow \alpha_{a_t} + 1$ . ▷ Increment number of “successes”.
- 7:     **else if**  $r_t == 0$  **then**  $\beta_{a_t} \leftarrow \beta_{a_t} + 1$ . ▷ Increment number of “failures”.

| Arm (i) | True $p_i$ | $\alpha_i$ | $\beta_i$ | $s_{i,t}$ |
|---------|------------|------------|-----------|-----------|
| 1       | 0.5        | 1          | 1         | 0.43      |
| 2       | 0.2        | 1          | 1         | 0.75      |
| 3       | 0.9        | 1          | 1         |           |

Sample from  $Beta(1,1)$  density →



| Time (t) | Arm Pulled ( $a_t$ ) | Reward ( $r_t$ ) |
|----------|----------------------|------------------|
| 1        |                      |                  |

# THOMPSON

## EXAMPLE

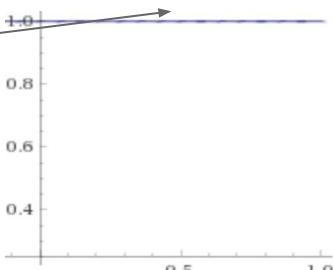


### Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ▷ Set  $Beta(\alpha_i, \beta_i)$  prior for each  $p_i$ .
- 2: **for**  $t = 1, 2, \dots, T$  **do**:
- 3:     For each arm  $i$ , get sample  $s_{i,t} \sim Beta(\alpha_i, \beta_i)$ . ▷ Each is a float in  $[0, 1]$ .
- 4:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} s_{i,t}$ . ▷ This “bakes in” exploration!
- 5:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 6:     **if**  $r_t == 1$  **then**  $\alpha_{a_t} \leftarrow \alpha_{a_t} + 1$ . ▷ Increment number of “successes”.
- 7:     **else if**  $r_t == 0$  **then**  $\beta_{a_t} \leftarrow \beta_{a_t} + 1$ . ▷ Increment number of “failures”.

| Arm (i) | True $p_i$ | $\alpha_i$ | $\beta_i$ | $s_{i,t}$ |
|---------|------------|------------|-----------|-----------|
| 1       | 0.5        | 1          | 1         | 0.43      |
| 2       | 0.2        | 1          | 1         | 0.75      |
| 3       | 0.9        | 1          | 1         | 0.11      |

Sample from  $Beta(1,1)$  density →



| Time (t) | Arm Pulled ( $a_t$ ) | Reward ( $r_t$ ) |
|----------|----------------------|------------------|
| 1        |                      |                  |

# THOMPSON

## EXAMPLE



### Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ▷ Set  $Beta(\alpha_i, \beta_i)$  prior for each  $p_i$ .
- 2: **for**  $t = 1, 2, \dots, T$  **do**:
- 3:     For each arm  $i$ , get sample  $s_{i,t} \sim Beta(\alpha_i, \beta_i)$ . ▷ Each is a float in  $[0, 1]$ .
- 4:     **Pull arm**  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} s_{i,t}$ . ▷ This “bakes in” exploration!
- 5:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 6:     **if**  $r_t == 1$  **then**  $\alpha_{a_t} \leftarrow \alpha_{a_t} + 1$ . ▷ Increment number of “successes”.
- 7:     **else if**  $r_t == 0$  **then**  $\beta_{a_t} \leftarrow \beta_{a_t} + 1$ . ▷ Increment number of “failures”.

| Arm (i) | True $p_i$ | $\alpha_i$ | $\beta_i$ | $s_{i,t}$ |
|---------|------------|------------|-----------|-----------|
| 1       | 0.5        | 1          | 1         | 0.43      |
| 2       | 0.2        | 1          | 1         | 0.75      |
| 3       | 0.9        | 1          | 1         | 0.11      |

| Time (t) | Arm Pulled ( $a_t$ ) | Reward ( $r_t$ ) |
|----------|----------------------|------------------|
| 1        | 2                    |                  |

Choose arm with highest sample!

# THOMPSON

## EXAMPLE



### Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ▷ Set  $Beta(\alpha_i, \beta_i)$  prior for each  $p_i$ .
- 2: **for**  $t = 1, 2, \dots, T$  **do**:
- 3:     For each arm  $i$ , get sample  $s_{i,t} \sim Beta(\alpha_i, \beta_i)$ . ▷ Each is a float in  $[0, 1]$ .
- 4:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} s_{i,t}$ . ▷ This “bakes in” exploration!
- 5:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 6:     **if**  $r_t == 1$  **then**  $\alpha_{a_t} \leftarrow \alpha_{a_t} + 1$ . ▷ Increment number of “successes”.
- 7:     **else if**  $r_t == 0$  **then**  $\beta_{a_t} \leftarrow \beta_{a_t} + 1$ . ▷ Increment number of “failures”.

| Arm (i) | True $p_i$ | $\alpha_i$ | $\beta_i$ | $s_{i,t}$ |
|---------|------------|------------|-----------|-----------|
| 1       | 0.5        | 1          | 1         | 0.43      |
| 2       | 0.2        | 1          | 1         | 0.75      |
| 3       | 0.9        | 1          | 1         | 0.11      |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 1        | 2                            | 0                        |

Observe reward 1 with probability 0.2  
and 0 with probability 0.8.

# THOMPSON

## EXAMPLE



### Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ▷ Set  $Beta(\alpha_i, \beta_i)$  prior for each  $p_i$ .
- 2: **for**  $t = 1, 2, \dots, T$  **do**:
- 3:     For each arm  $i$ , get sample  $s_{i,t} \sim Beta(\alpha_i, \beta_i)$ . ▷ Each is a float in  $[0, 1]$ .
- 4:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} s_{i,t}$ . ▷ This “bakes in” exploration!
- 5:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 6:     **if**  $r_t == 1$  **then**  $\alpha_{a_t} \leftarrow \alpha_{a_t} + 1$ . ▷ Increment number of “successes”.
- 7:     **else if**  $r_t == 0$  **then**  $\beta_{a_t} \leftarrow \beta_{a_t} + 1$ . ▷ Increment number of “failures”.

| Arm (i) | True $p_i$ | $\alpha_i$ | $\beta_i$ | $s_{i,t}$ |
|---------|------------|------------|-----------|-----------|
| 1       | 0.5        | 1          | 1         |           |
| 2       | 0.2        | 1          | 2         |           |
| 3       | 0.9        | 1          | 1         |           |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 1        | 2                            | 0                        |

Add a count of 1 to the failures :(.

# THOMPSON

## EXAMPLE

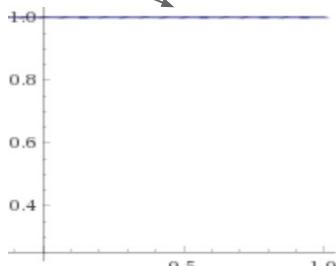


### Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ▷ Set  $Beta(\alpha_i, \beta_i)$  prior for each  $p_i$ .
- 2: **for**  $t = 1, 2, \dots, T$  **do**:
- 3:     For each arm  $i$ , get sample  $s_{i,t} \sim Beta(\alpha_i, \beta_i)$ . ▷ Each is a float in  $[0, 1]$ .
- 4:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} s_{i,t}$ . ▷ This “bakes in” exploration!
- 5:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 6:     **if**  $r_t == 1$  **then**  $\alpha_{a_t} \leftarrow \alpha_{a_t} + 1$ . ▷ Increment number of “successes”.
- 7:     **else if**  $r_t == 0$  **then**  $\beta_{a_t} \leftarrow \beta_{a_t} + 1$ . ▷ Increment number of “failures”.

| Arm (i) | True $p_i$ | $\alpha_i$ | $\beta_i$ | $s_{i,t}$ |
|---------|------------|------------|-----------|-----------|
| 1       | 0.5        | 1          | 1         | 0.52      |
| 2       | 0.2        | 1          | 2         |           |
| 3       | 0.9        | 1          | 1         |           |

Sample from  $Beta(1,1)$  density →



| Time (t) | Arm Pulled ( $a_t$ ) | Reward ( $r_t$ ) |
|----------|----------------------|------------------|
| 2        |                      |                  |

# THOMPSON

## EXAMPLE

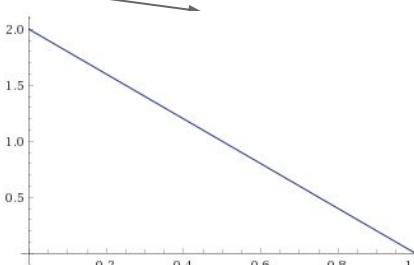


### Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ▷ Set  $Beta(\alpha_i, \beta_i)$  prior for each  $p_i$ .
- 2: **for**  $t = 1, 2, \dots, T$  **do**:
- 3:     For each arm  $i$ , get sample  $s_{i,t} \sim Beta(\alpha_i, \beta_i)$ . ▷ Each is a float in  $[0, 1]$ .
- 4:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} s_{i,t}$ . ▷ This “bakes in” exploration!
- 5:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 6:     **if**  $r_t == 1$  **then**  $\alpha_{a_t} \leftarrow \alpha_{a_t} + 1$ . ▷ Increment number of “successes”.
- 7:     **else if**  $r_t == 0$  **then**  $\beta_{a_t} \leftarrow \beta_{a_t} + 1$ . ▷ Increment number of “failures”.

| Arm (i) | True $p_i$ | $\alpha_i$ | $\beta_i$ | $s_{i,t}$ |
|---------|------------|------------|-----------|-----------|
| 1       | 0.5        | 1          | 1         | 0.52      |
| 2       | 0.2        | 1          | 2         | 0.05      |
| 3       | 0.9        | 1          | 1         |           |

Sample from  $Beta(1,2)$  density →



| Time (t) | Arm Pulled ( $a_t$ ) | Reward ( $r_t$ ) |
|----------|----------------------|------------------|
| 2        |                      |                  |

# THOMPSON

## EXAMPLE

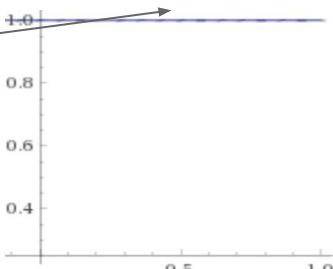


### Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ▷ Set  $Beta(\alpha_i, \beta_i)$  prior for each  $p_i$ .
- 2: **for**  $t = 1, 2, \dots, T$  **do**:
- 3:     For each arm  $i$ , get sample  $s_{i,t} \sim Beta(\alpha_i, \beta_i)$ . ▷ Each is a float in  $[0, 1]$ .
- 4:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} s_{i,t}$ . ▷ This “bakes in” exploration!
- 5:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 6:     **if**  $r_t == 1$  **then**  $\alpha_{a_t} \leftarrow \alpha_{a_t} + 1$ . ▷ Increment number of “successes”.
- 7:     **else if**  $r_t == 0$  **then**  $\beta_{a_t} \leftarrow \beta_{a_t} + 1$ . ▷ Increment number of “failures”.

| Arm (i) | True $p_i$ | $\alpha_i$ | $\beta_i$ | $s_{i,t}$ |
|---------|------------|------------|-----------|-----------|
| 1       | 0.5        | 1          | 1         | 0.52      |
| 2       | 0.2        | 1          | 2         | 0.05      |
| 3       | 0.9        | 1          | 1         | 0.67      |

Sample from  $Beta(1,1)$  density →



| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 2        |                              |                          |

# THOMPSON

## EXAMPLE



### Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ▷ Set  $Beta(\alpha_i, \beta_i)$  prior for each  $p_i$ .
- 2: **for**  $t = 1, 2, \dots, T$  **do**:
- 3:     For each arm  $i$ , get sample  $s_{i,t} \sim Beta(\alpha_i, \beta_i)$ . ▷ Each is a float in  $[0, 1]$ .
- 4:     **Pull arm**  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} s_{i,t}$ . ▷ This “bakes in” exploration!
- 5:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 6:     **if**  $r_t == 1$  **then**  $\alpha_{a_t} \leftarrow \alpha_{a_t} + 1$ . ▷ Increment number of “successes”.
- 7:     **else if**  $r_t == 0$  **then**  $\beta_{a_t} \leftarrow \beta_{a_t} + 1$ . ▷ Increment number of “failures”.

| Arm (i) | True $p_i$ | $\alpha_i$ | $\beta_i$ | $s_{i,t}$ |
|---------|------------|------------|-----------|-----------|
| 1       | 0.5        | 1          | 1         | 0.52      |
| 2       | 0.2        | 1          | 2         | 0.05      |
| 3       | 0.9        | 1          | 1         | 0.67      |

| Time (t) | Arm Pulled ( $a_t$ ) | Reward ( $r_t$ ) |
|----------|----------------------|------------------|
| 2        | 3                    |                  |

Choose arm with highest sample!

# THOMPSON

## EXAMPLE



### Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ▷ Set  $Beta(\alpha_i, \beta_i)$  prior for each  $p_i$ .
- 2: **for**  $t = 1, 2, \dots, T$  **do**:
- 3:     For each arm  $i$ , get sample  $s_{i,t} \sim Beta(\alpha_i, \beta_i)$ . ▷ Each is a float in  $[0, 1]$ .
- 4:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} s_{i,t}$ . ▷ This “bakes in” exploration!
- 5:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 6:     **if**  $r_t == 1$  **then**  $\alpha_{a_t} \leftarrow \alpha_{a_t} + 1$ . ▷ Increment number of “successes”.
- 7:     **else if**  $r_t == 0$  **then**  $\beta_{a_t} \leftarrow \beta_{a_t} + 1$ . ▷ Increment number of “failures”.

| Arm (i) | True $p_i$ | $\alpha_i$ | $\beta_i$ | $s_{i,t}$ |
|---------|------------|------------|-----------|-----------|
| 1       | 0.5        | 1          | 1         | 0.52      |
| 2       | 0.2        | 1          | 2         | 0.05      |
| 3       | 0.9        | 1          | 1         | 0.67      |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 2        | 3                            | 1                        |

Observe reward 1 with probability 0.9  
and 0 with probability 0.1.

# THOMPSON

## EXAMPLE



### Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ▷ Set  $Beta(\alpha_i, \beta_i)$  prior for each  $p_i$ .
- 2: **for**  $t = 1, 2, \dots, T$  **do**:
- 3:     For each arm  $i$ , get sample  $s_{i,t} \sim Beta(\alpha_i, \beta_i)$ . ▷ Each is a float in  $[0, 1]$ .
- 4:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} s_{i,t}$ . ▷ This “bakes in” exploration!
- 5:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 6:     **if**  $r_t == 1$  **then**  $\alpha_{a_t} \leftarrow \alpha_{a_t} + 1$ . ▷ Increment number of “successes”.
- 7:     **else if**  $r_t == 0$  **then**  $\beta_{a_t} \leftarrow \beta_{a_t} + 1$ . ▷ Increment number of “failures”.

| Arm (i) | True $p_i$ | $\alpha_i$ | $\beta_i$ | $s_{i,t}$ |
|---------|------------|------------|-----------|-----------|
| 1       | 0.5        | 1          | 1         |           |
| 2       | 0.2        | 1          | 2         |           |
| 3       | 0.9        | 2          | 1         |           |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 2        | 3                            | 1                        |

Add a count of 1 to the successes :).

# THOMPSON

## EXAMPLE

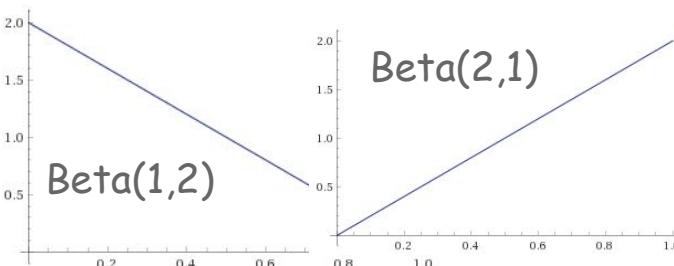
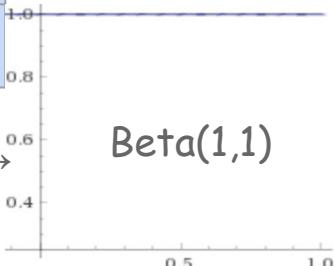


### Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ▷ Set  $Beta(\alpha_i, \beta_i)$  prior for each  $p_i$ .
- 2: **for**  $t = 1, 2, \dots, T$  **do**:
- 3:     For each arm  $i$ , get sample  $s_{i,t} \sim Beta(\alpha_i, \beta_i)$ . ▷ Each is a float in  $[0, 1]$ .
- 4:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} s_{i,t}$ . ▷ This “bakes in” exploration!
- 5:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 6:     **if**  $r_t == 1$  **then**  $\alpha_{a_t} \leftarrow \alpha_{a_t} + 1$ . ▷ Increment number of “successes”.
- 7:     **else if**  $r_t == 0$  **then**  $\beta_{a_t} \leftarrow \beta_{a_t} + 1$ . ▷ Increment number of “failures”.

| Arm (i) | True $p_i$ | $\alpha_i$ | $\beta_i$ | $s_{i,t}$ |
|---------|------------|------------|-----------|-----------|
| 1       | 0.5        | 1          | 1         | 0.44      |
| 2       | 0.2        | 1          | 2         | 0.27      |
| 3       | 0.9        | 2          | 1         | 0.86      |

Sample from each arm's Beta distribution →



| Time (t) | Arm Pulled ( $a_t$ ) | Reward ( $r_t$ ) |
|----------|----------------------|------------------|
| 3        |                      |                  |

Beta(2,1)

# THOMPSON

## EXAMPLE



### Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ▷ Set  $Beta(\alpha_i, \beta_i)$  prior for each  $p_i$ .
- 2: **for**  $t = 1, 2, \dots, T$  **do**:
- 3:     For each arm  $i$ , get sample  $s_{i,t} \sim Beta(\alpha_i, \beta_i)$ . ▷ Each is a float in  $[0, 1]$ .
- 4:     **Pull arm**  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} s_{i,t}$ . ▷ This “bakes in” exploration!
- 5:     **Receive reward**  $r_t \sim Ber(p_{a_t})$ .
- 6:     **if**  $r_t == 1$  **then**  $\alpha_{a_t} \leftarrow \alpha_{a_t} + 1$ . ▷ Increment number of “successes”.
- 7:     **else if**  $r_t == 0$  **then**  $\beta_{a_t} \leftarrow \beta_{a_t} + 1$ . ▷ Increment number of “failures”.

| Arm (i) | True $p_i$ | $\alpha_i$ | $\beta_i$ | $s_{i,t}$ |
|---------|------------|------------|-----------|-----------|
| 1       | 0.5        | 1          | 1         | 0.44      |
| 2       | 0.2        | 1          | 2         | 0.27      |
| 3       | 0.9        | 2          | 1         | 0.86      |

| Time (t) | Arm Pulled ( $a_t$ ) | Reward ( $r_t$ ) |
|----------|----------------------|------------------|
| 3        | 3                    | 0                |

Observe reward 1 with probability 0.9  
and 0 with probability 0.1.

# THOMPSON

## EXAMPLE



### Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ▷ Set  $Beta(\alpha_i, \beta_i)$  prior for each  $p_i$ .
- 2: **for**  $t = 1, 2, \dots, T$  **do**:
- 3:     For each arm  $i$ , get sample  $s_{i,t} \sim Beta(\alpha_i, \beta_i)$ . ▷ Each is a float in  $[0, 1]$ .
- 4:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} s_{i,t}$ . ▷ This “bakes in” exploration!
- 5:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 6:     **if**  $r_t == 1$  **then**  $\alpha_{a_t} \leftarrow \alpha_{a_t} + 1$ . ▷ Increment number of “successes”.
- 7:     **else if**  $r_t == 0$  **then**  $\beta_{a_t} \leftarrow \beta_{a_t} + 1$ . ▷ Increment number of “failures”.

| Arm (i) | True $p_i$ | $\alpha_i$ | $\beta_i$ | $s_{i,t}$ |
|---------|------------|------------|-----------|-----------|
| 1       | 0.5        | 1          | 1         | 0.44      |
| 2       | 0.2        | 1          | 2         | 0.27      |
| 3       | 0.9        | 2          | 2         | 0.86      |

| Time (t) | Arm Pulled (a <sub>t</sub> ) | Reward (r <sub>t</sub> ) |
|----------|------------------------------|--------------------------|
| 3        | 3                            | 0                        |

Add a count of 1 to the failures :(.

# THOMPSON

## EXAMPLE

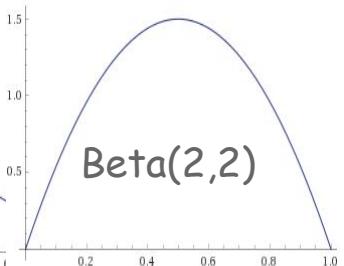
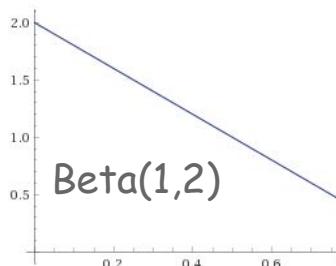
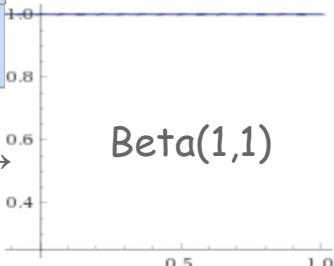


### Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ▷ Set  $Beta(\alpha_i, \beta_i)$  prior for each  $p_i$ .
- 2: **for**  $t = 1, 2, \dots, T$  **do**:
- 3:     For each arm  $i$ , get sample  $s_{i,t} \sim Beta(\alpha_i, \beta_i)$ . ▷ Each is a float in  $[0, 1]$ .
- 4:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} s_{i,t}$ . ▷ This “bakes in” exploration!
- 5:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 6:     **if**  $r_t == 1$  **then**  $\alpha_{a_t} \leftarrow \alpha_{a_t} + 1$ . ▷ Increment number of “successes”.
- 7:     **else if**  $r_t == 0$  **then**  $\beta_{a_t} \leftarrow \beta_{a_t} + 1$ . ▷ Increment number of “failures”.

| Arm (i) | True $p_i$ | $\alpha_i$ | $\beta_i$ | $s_{i,t}$ |
|---------|------------|------------|-----------|-----------|
| 1       | 0.5        | 1          | 1         | 0.63      |
| 2       | 0.2        | 1          | 2         | 0.15      |
| 3       | 0.9        | 2          | 2         | 0.44      |

Sample from each arm's Beta distribution →



| Time (t) | Arm Pulled ( $a_t$ ) | Reward ( $r_t$ ) |
|----------|----------------------|------------------|
| 4        |                      |                  |

# THOMPSON

## EXAMPLE



### Algorithm 5 Thompson Sampling Algorithm for Beta-Bernoulli Bandits

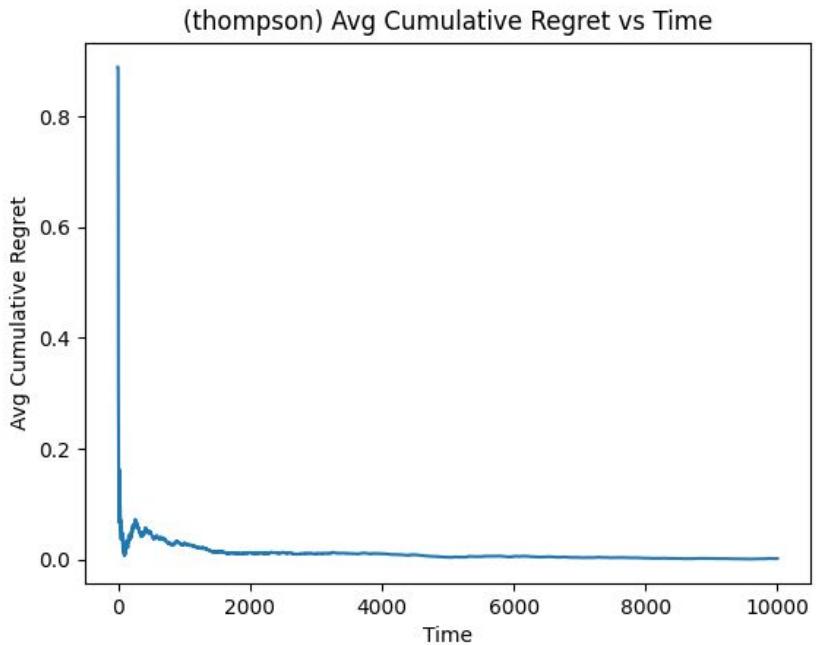
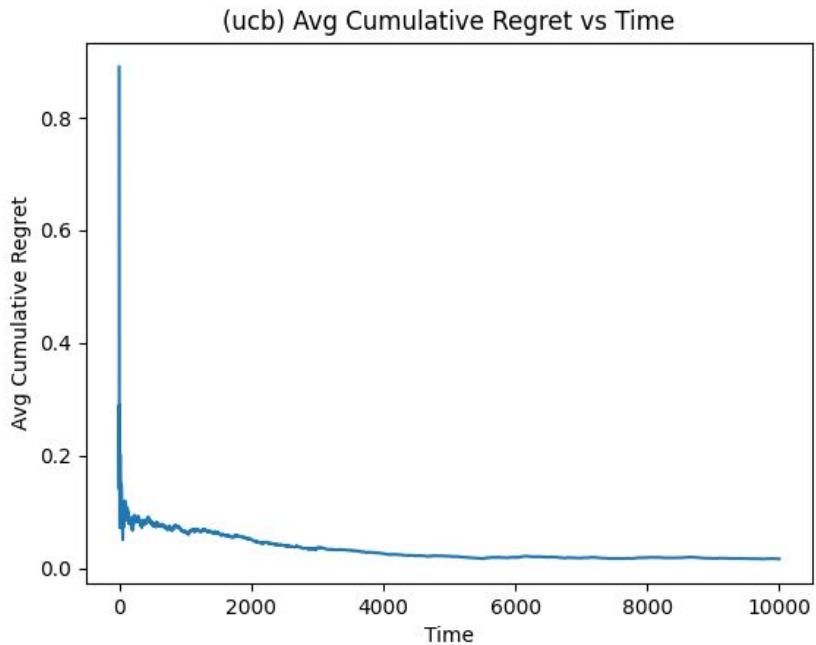
- 1: For each arm  $i \in \{1, \dots, K\}$ , initialize  $\alpha_i = \beta_i = 1$ . ▷ Set  $Beta(\alpha_i, \beta_i)$  prior for each  $p_i$ .
- 2: **for**  $t = 1, 2, \dots, T$  **do**:
- 3:     For each arm  $i$ , get sample  $s_{i,t} \sim Beta(\alpha_i, \beta_i)$ . ▷ Each is a float in  $[0, 1]$ .
- 4:     Pull arm  $a_t = \arg \max_{i \in \{1, 2, \dots, K\}} s_{i,t}$ . ▷ This “bakes in” exploration!
- 5:     Receive reward  $r_t \sim Ber(p_{a_t})$ .
- 6:     **if**  $r_t == 1$  **then**  $\alpha_{a_t} \leftarrow \alpha_{a_t} + 1$ . ▷ Increment number of “successes”.
- 7:     **else if**  $r_t == 0$  **then**  $\beta_{a_t} \leftarrow \beta_{a_t} + 1$ . ▷ Increment number of “failures”.

| Arm (i) | True $p_i$ | $\alpha_i$ | $\beta_i$ | $s_{i,t}$ |
|---------|------------|------------|-----------|-----------|
| 1       | 0.5        | 1          | 1         | 0.63      |
| 2       | 0.2        | 1          | 2         | 0.15      |
| 3       | 0.9        | 2          | 2         | 0.44      |

| Time (t) | Arm Pulled ( $a_t$ ) | Reward ( $r_t$ ) |
|----------|----------------------|------------------|
| 4        |                      |                  |

And so on!!! Notice how we explore because there's some chance the “best” arm will have a lower sample occasionally and let other arms win!

# UCB VS THOMPSON SAMPLING: AVG REGRET OVER TIME



# UCB VS THOMPSON SAMPLING: PROPORTION OF TIMES PULLED

