

Reinforcement Learning for Flappy Bird using Tabular Q-Learning

AMJAD K P
IMSC CS AI & DS

November 6, 2025

Abstract

This report details the implementation of a Tabular Q-Learning agent designed to play the game Flappy Bird. The primary challenge lies in the game's continuous state space, which is incompatible with a standard Q-table. We address this by discretizing the state into a finite grid. This paper outlines the environment, the state-action-reward design, the discretization method, and the Q-Learning algorithm. We analyze the agent's performance by visualizing its learning curve and the final learned policy (Q-table), demonstrating that this simple approach is sufficient to learn a competent policy.

Contents

| | | |
|----------|---|----------|
| 1 | Introduction | 2 |
| 2 | Methodology | 2 |
| 2.1 | Experimental Environment | 2 |
| 2.2 | State Discretization | 2 |
| 2.3 | Algorithm: Q-Learning | 3 |
| 2.4 | Hyperparameters | 3 |
| 3 | Experimental Results and Analysis | 3 |
| 3.1 | Learning Curve | 3 |
| 3.2 | Q-Table Policy Visualization | 4 |
| 4 | Conclusion | 5 |
| A | Complete Training and Visualization Code | 6 |

1 Introduction

Reinforcement Learning (RL) is a paradigm of machine learning where an agent learns to make optimal decisions by interacting with an environment. The agent receives rewards or penalties for its actions, with the goal of maximizing its cumulative reward. Q-Learning is a popular model-free, off-policy RL algorithm that learns a quality function (Q-value) for each state-action pair.

The game Flappy Bird presents a classic RL problem. The agent (the bird) must learn a policy—when to flap or not—to navigate through a series of pipes. The game’s challenge stems from its continuous state space; the bird’s position and velocity are floating-point numbers, leading to a theoretically infinite number of states. This project tackles this by discretizing the state space to make it solvable with a traditional Q-table.

2 Methodology

2.1 Experimental Environment

The project uses the `flappy-bird-gymnasium` environment, a standard interface for the Flappy Bird game that follows the Gymnasium API.

- **Actions (A):** The agent has two discrete actions:
 - 0: Do nothing (let gravity act)
 - 1: Flap (apply upward velocity)
- **Reward (R):** We implemented a custom reward structure (reward engineering) to guide the agent:
 - +0.1 for every frame the bird stays alive.
 - -100 (large penalty) for crashing (game over).
- **State (S):** The raw state provided by the environment consists of two continuous values:
 - `horizontal_distance`: The bird’s X-distance to the next pipe.
 - `vertical_distance`: The bird’s Y-distance to the center of the next pipe’s gap.

2.2 State Discretization

To use a Q-table, we must convert the continuous state into a finite number of discrete states. We defined a grid of (20, 20) bins. The continuous horizontal and vertical distances were “bucketed” into one of these 20 bins each.

This transforms a continuous state like (14.123, -5.456) into a discrete tuple like (3, 8), which can be used as a key in our Q-table. This results in a total of $20 \times 20 = 400$ possible states.

2.3 Algorithm: Q-Learning

We used the standard Q-Learning (Temporal-Difference) update rule. After taking an action a in state s and observing the reward r and next state s' , the Q-value is updated using the Bellman equation:

$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)] \quad (1)$$

Where α is the learning rate and γ is the discount factor.

2.4 Hyperparameters

The agent was trained using the parameters shown in Table 1. An ϵ -greedy strategy was used for exploration, where ϵ decayed over time from 1.0 to 0.01.

Table 1: Hyperparameters for Q-Learning Training

| Parameter | Value |
|------------------------------|----------|
| Training Episodes | 50,000 |
| Learning Rate (α) | 0.1 |
| Discount Factor (γ) | 0.99 |
| State Bins | (20, 20) |
| Epsilon Start | 1.0 |
| Epsilon End | 0.01 |
| Epsilon Decay Rate | 0.99995 |

3 Experimental Results and Analysis

3.1 Learning Curve

The agent's performance was tracked by logging the total reward for each episode. Figure 1 shows a moving average of the episodic rewards. There is a clear upward trend, indicating that the agent successfully learned a policy to increase its survival time and, therefore, its cumulative reward. The initial high variance is due to random exploration, which settles as the agent begins to exploit its learned policy.




Figure 1: Agent learning curve (moving average of 1000 episodes).

3.2 Q-Table Policy Visualization

To understand *what* the agent learned, we visualized the Q-table as a heatmap in Figure 2. The color of each cell (h, v) represents the maximum Q-value (the expected future reward) for that discretized state.

A clear "path" of high-value (yellow) states is visible. This path represents the optimal position (relative to the pipe) that the agent learned to stay in. The surrounding blue areas represent low-value states, such as being too high or too low, which the agent learned are dangerous and lead to a crash. This visual confirms that the agent learned a coherent and optimal policy.




Figure 2: Heatmap of the max Q-value for each discretized state.

4 Conclusion

This project successfully demonstrated that a simple Tabular Q-Learning method, when combined with careful state discretization, is capable of learning a competent policy for the Flappy Bird game. The learning curve shows clear improvement over time, and the Q-table heatmap provides a clear visualization of the optimal "path" the agent discovered.

While effective, this method is sensitive to the number of bins chosen. Future work could explore Deep Q-Networks (DQN), which use a neural network to approximate the Q-function, thereby eliminating the need for manual state discretization and allowing the agent to learn from the raw, continuous state.

A Complete Training and Visualization Code

```
1 import gymnasium as gym
2 import flappy_bird_gymnasium
3 import numpy as np
4 import time
5 import pickle
6 import matplotlib.pyplot as plt
7
8 # --- Step 1 & 2: Environment Setup and Discretization ---
9
10 print(" --- Step 1 & 2: Setting up Environment and Discretization ---")
11
12 # Create the environment to check its observation space
13 env = gym.make("FlappyBird-v0")
14
15 # Define how many "buckets" we want for each part of the state
16 # (horizontal_bins, vertical_bins)
17 STATE_BINS = (20, 20)
18
19 # Get the min and max values for each state variable
20 state_low = env.observation_space.low
21 state_high = env.observation_space.high
22
23 # Create the bin edges for digitizing
24 horizontal_bins = np.linspace(state_low[0], state_high[0], STATE_BINS[0] - 1)
25 vertical_bins = np.linspace(state_low[1], state_high[1], STATE_BINS[1] - 1)
26
27 def discretize_state(state):
28     """
29         Converts a continuous state (like [14.123, -5.456])
30         into a discrete tuple (like (3, 8)).
31     """
32     horizontal_bin = np.digitize(state[0], horizontal_bins)
33     vertical_bin = np.digitize(state[1], vertical_bins)
34
35     return (horizontal_bin, vertical_bin)
36
37 # --- Test discretization ---
38 test_state_continuous, info = env.reset()
39 test_state_discrete = discretize_state(test_state_continuous)
40 print(f"Continuous state example: {test_state_continuous}")
41 print(f"Discrete state example: {test_state_discrete}")
42 env.close()
43
44 # --- Step 3: Set Up Hyperparameters and Q-Table ---
45
46 print(" --- Step 3: Initializing Hyperparameters and Q-Table ---")
47
48 # Create our Q-table as a dictionary
49 q_table = {}
50
51 # --- Hyperparameters ---
52 EPISODES = 50_000          # Total number of games to play
53 LEARNING_RATE = 0.1        # (alpha)
54 DISCOUNT_FACTOR = 0.99    # (gamma)
```

```

55
56 # --- Epsilon-Greedy Strategy ---
57 EPSILON_START = 1.0          # Start with 100% random actions
58 EPSILON_END = 0.01           # End with 1% random actions
59 EPSILON_DECAY = 0.99995     # How fast epsilon shrinks
60 epsilon = EPSILON_START
61
62 # How many actions are there? (0=do_nothing, 1=flap)
63 ACTION_COUNT = env.action_space.n
64
65 # --- For plotting ---
66 episode_rewards = []
67
68 # --- Step 4: The Main Training Loop ---
69
70 print("--- Step 4: Starting Training Loop ---")
71
72 # This environment is just for training (no rendering)
73 env_train = gym.make("FlappyBird-v0")
74
75 for episode in range(EPSISODES):
76
77     # Start a new game
78     state_continuous, info = env_train.reset()
79     state = discretize_state(state_continuous)
80
81     total_reward = 0 # Track reward for this episode
82     done = False
83
84     while not done:
85
86         # 1. Epsilon-Greedy: Choose an action
87         if np.random.uniform(0, 1) < epsilon:
88             action = env_train.action_space.sample()
89         else:
90             action = np.argmax(q_table.get(state, np.zeros(ACTION_COUNT
91             )))
92
93         # 2. Take the action and get feedback
94         next_state_continuous, reward, done, truncated, info =
95         env_train.step(action)
96         next_state = discretize_state(next_state_continuous)
97
98         # 3. Reward Engineering
99         if done:
100             reward = -100 # Big penalty for dying
101
102             total_reward += reward # Log reward
103
104             # 4. The Q-Learning Update Rule
105             current_q_values = q_table.get(state, np.zeros(ACTION_COUNT))
106             max_future_q = np.max(q_table.get(next_state, np.zeros(
107             ACTION_COUNT)))
108             current_q = current_q_values[action]
109
110             new_q = current_q + LEARNING_RATE * (reward + DISCOUNT_FACTOR *
111             max_future_q - current_q)

```

```

109     # 5. Update Q-table
110     if state not in q_table:
111         q_table[state] = np.zeros(ACTION_COUNT)
112
113     q_table[state][action] = new_q
114
115     # 6. Update state
116     state = next_state
117
118     # --- End of episode ---
119
120     # Decay epsilon
121     epsilon = max(EPSILON_END, epsilon * EPSILON_DECAY)
122
123     # Log the total reward for plotting
124     episode_rewards.append(total_reward)
125
126     # Print progress
127     if (episode + 1) % 5000 == 0:
128         print(f"Episode: {episode + 1}/{EPISODES} | Epsilon: {epsilon :.4f}")
129
130 env_train.close()
131 print("Training finished!")
132
133 # --- Step 5: Visualize Learning Curve ---
134
135 print("--- Step 5: Generating Learning Curve Plot ---")
136
137 # Calculate a moving average
138 chunk_size = 1000
139 moving_averages = [
140     np.mean(episode_rewards[i:i + chunk_size])
141     for i in range(0, len(episode_rewards), chunk_size)
142 ]
143
144 plt.figure(figsize=(10, 6))
145 plt.plot(
146     np.arange(len(moving_averages)) * chunk_size,
147     moving_averages,
148     label=f'Moving Average ({chunk_size} episodes)',
149 )
150 plt.title("Average Reward per Episode over Time")
151 plt.xlabel("Episode")
152 plt.ylabel("Average Reward")
153 plt.legend()
154 plt.grid(True)
155
156 # Save the plot to a file
157 plt.savefig("flappy_learning_curve.png")
158 print("Learning curve saved to 'flappy_learning_curve.png'")
159 # plt.show() # Uncomment this if you want to see the plot immediately
160
161 # --- Step 6: Visualize the Q-Table (Heatmap) ---
162
163 print("--- Step 6: Generating Q-Table Heatmap ---")
164
165 def get_max_q_value(state):

```

```

166     """Returns the max Q-value for a given state, or 0 if state is
167     unknown."""
168     return np.max(q_table.get(state, np.zeros(ACTION_COUNT)))
169
170 # Create a 2D grid representing our state space
171 heatmap_data = np.zeros(STATE_BINS)
172
173 for h_bin in range(STATE_BINS[0]):
174     for v_bin in range(STATE_BINS[1]):
175         state = (h_bin, v_bin)
176         heatmap_data[v_bin, h_bin] = get_max_q_value(state) # v, h
177         order for plotting
178
179 # Create the heatmap
180 plt.figure(figsize=(10, 8))
181 plt.imshow(heatmap_data, cmap='viridis', interpolation='nearest',
182             aspect='auto')
183
184 plt.title("Q-Table Value Heatmap")
185 plt.xlabel("Horizontal Distance to Pipe (Discretized)")
186 plt.ylabel("Vertical Distance to Pipe (Discretized)")
187 plt.colorbar(label="Max Q-Value (Expected Future Reward)")
188
189 # Save the plot
190 plt.savefig("flappy-q_table_heatmap.png")
191 print("Q-table heatmap saved to 'flappy-q_table_heatmap.png'")
192 # plt.show() # Uncomment to show
193
194 # --- Save the Q-Table ---
195 print("--- Saving Q-Table to pkl file ---")
196 with open("flappy-q_table.pkl", "wb") as f:
197     pickle.dump(q_table, f)
198
199 print(f"Q-table saved! It has {len(q_table)} states.")
200
201 # --- Step 7: See Your Trained Agent in Action! ---
202
203 print("--- Step 7: Running Trained Agent (Human View) ---")
204 env_human = gym.make("FlappyBird-v0", render_mode="human")
205
206 for episode in range(5):
207     state_continuous, info = env_human.reset()
208     state = discretize_state(state_continuous)
209
210     done = False
211
212     while not done:
213         # We ALWAYS exploit (no epsilon)
214         action = np.argmax(q_table.get(state, np.zeros(ACTION_COUNT)))
215
216         # Take the best action
217         next_state_continuous, reward, done, truncated, info =
218         env_human.step(action)
219         state = discretize_state(next_state_continuous)
220
221         # Slow it down so we can watch
222         time.sleep(1 / 60) # 60 FPS

```

```
220  
221 env_human.close()  
222 print("Agent run finished.")
```

Listing 1: flappy.py: Full Training Script